

AIR:

An IR for PIF

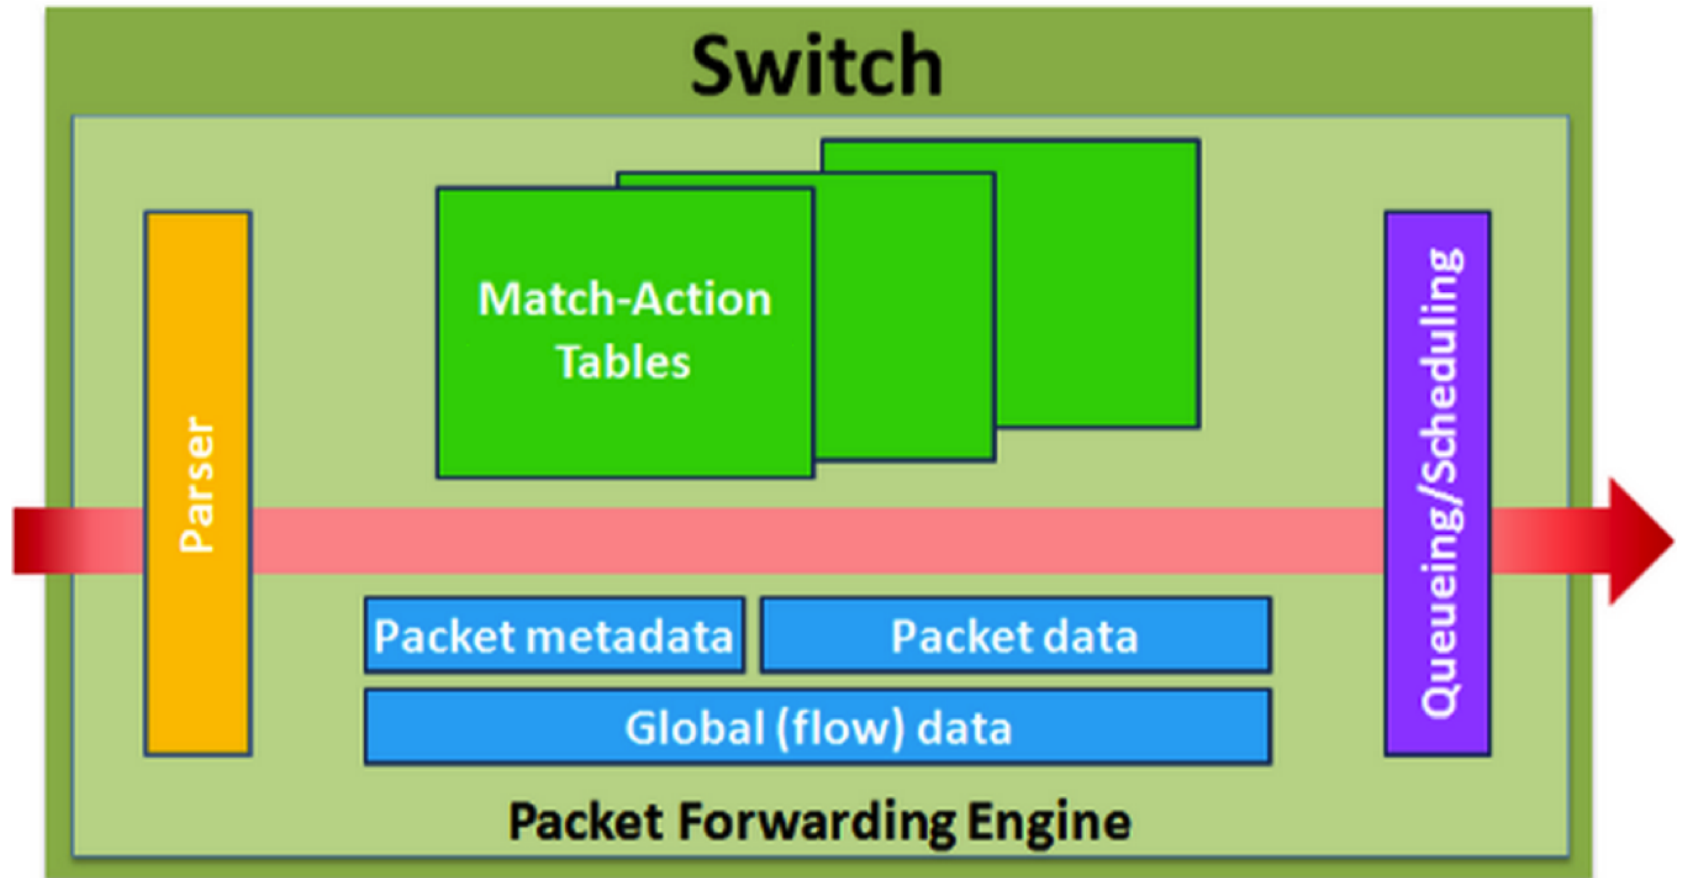
A proposed Intermediate Representation for
Protocol Independent Forwarding

Dan Talayco
September, 2014

IR Overview

- Thin waist between
 - Multiple high level languages
 - Multiple target devices
- Driven by Abstract Forwarding Model
 - Should capture all semantics of that model
- Extensible and minimal

The Abstract Forwarding Model




What's in an Instance?

- Parse Graph
- Control Flow Graph
- Table Dependency Graph (maybe)
- Tables
- Fields and Headers
- Actions

Graphs are the dominant term

Hmm, what's good for graphs...
dot?

```
digraph G {  
    start -> ethernet;  
    ethernet -> vlan [label="0x8100"];  
    ethernet -> ip4 [label="0x0800"];  
    vlan -> ip4 [label="0x0800"];  
    ip4 -> udp [label=6];  
    ip4 -> tcp [label=17];  
}
```



dot dot dot

- Many tools for dot
 - Visualization tools
 - Analysis and processing libraries import dot
- But...
 - We don't care about many dot attributes
 - We may want to add our own attributes
- So...
 - Assume we can update dot tools to support (or at least ignore) the attributes we want to add
 - More in a minute...

Syntax Proposal

- **YAML for declarations and object attributes**
 - But the metalanguage isn't that important
 - Can express these ideas in anything reasonable
- **Define a scheme for legal IR instances**
 - Build a validator as a first step
- **Components**
 - Tables, Fields, Actions, Parse Graph, CFG, TDG
- **Graph representation is selectable**
 - `format : <supported-format>`
 - Graph itself is a blob in the proper format
 - Will show with dot today

A Simple Example (excerpts)

Header declaration

```
ethernet :  
  type : header  
  doc : "The L2 header"  
  fields :  
    - dst_mac : 48  
    - src_mac : 48  
    - ethertype : 16
```

Parser State

```
ethernet_p :  
  type : parse_state  
  extracts :  
    - ethernet  
  select_value :  
    - ethernet.ethertype
```

Table declaration

```
13 :  
  type : table  
  doc : "Do L3 intf selection"  
  match_on :  
    metadata.vfi : exact  
    ipv4.dst : exact
```

Continued

Example, continued

Parse Graph

```
parser :  
  type : parse_graph  
  doc  : "Implementation of primary parser"  
  format : dot  
  implementation : >  
    start -> ethernet_p  
    ethernet_p -> vlan_p [value=0x8100]  
    ethernet_p -> ip4_p [value=0x0800]  
    vlan_p -> ip4_p [value=0x0800]  
    ip4_p -> udp_p [value=6]  
    ip4_p -> tcp_p [value=17]
```

Continued

Example, continued

Control flow

```
ingress_flow :  
  type : control_flow_graph  
  format : dot  
  # Use "action" here; really just a label  
  implementation : >  
    vlan -> l2 [action="set_l2_vfi_a"]  
    vlan -> l3 [action="set_l3_vfi_a"]  
    l2 -> flood_block [action="flood"]  
    l2 -> queue [action="set_egress_spec"]  
    l3 -> apply_route [action="set_route"]  
    apply_route -> queue [action="apply_route_if"]
```

Also part of the IR

- **Fields**
 - Identify and define attributes
 - Support extended attributes
- **Actions**
 - The list of primitive actions
 - An explanation of their semantics
 - Mechanism to extend the list...but shouldn't encourage
 - Define semantics for combining primitive actions into actions used by tables
- **Table matching semantics**
 - Define: exact, ternary, etc

Details of AIR

- Instance driven declaration
- Each AIR object has explicit type specified
- Object types shown on next slide
 - field_type, header, etc
- All AIR objects must have 'type' and 'doc' attributes
- May use YAML mechanisms such as references

AIR object types and their supported attributes

- **field_type**
 - width
 - signed
 - saturating
 - ...
- **header**
 - fields (list)
- **parser_state**
 - extracts
 - select_value
- **parse_graph**
 - format
 - implementation

Continued

AIR object types, continued

- **action**
 - format
 - parameter_list
 - each w/ attrs
 - implementation
- **table**
 - match_on
- **control_flow_graph**
 - format
 - implementation
- **dependency_graph**
 - format
 - implementation

Proposed Timeline

- Validator for AIR
 - Q4 2014
- Simulator for AIR
 - Q1 2015
- AIR tools
 - Visualizer, editors, formal model checking
 - Starting in Q4 2014
- AIR for OVS
 - Q1/Q2 2015

Questions

- Binding Time
 - Should AIR initially support run-time binding of field names for action primitives?
 - Proposal: No.
 - Reduces (removes?) need for TDG
 - Consider support in the future
- Options for Semantics for Action Primitives
 - Documentation
 - Optionally with reference implementation
 - Define underlying compute model and give algos

Tasks and Sign-up Sheet

- Details for AIR-YAML specification
- Complete "basic functionality" example
- Simulations
 - Python
 - HW Simulation Environment driven
 - (more) performant C-based model
- OVS Integration
- Visualization Tools (packet tracker?)
- Devise Debugging and Monitoring tools
- Interactive Designer