



Some IR Questions

Ben Mack-Crane

Huawei

What is required in an IR?



- Match Action Table definitions
- Data definitions
 - Packet data (Header)
 - Packet metadata
 - Global and/or flow table metadata
- Action definitions
 - Invokable procedures that operate on packets and metadata
- Event definitions?
 - Are these variable or can we define a fixed set of event types?
 - Packet received
 - Timer expired
 - Condition met (e.g., threshold crossing)



What are we experimenting with?



- Parser
- Control Flow
- Hardware Abstraction

How is this done with OpenFlow datapath models today?

What can we add in the IR and what are the tradeoffs?



OpenFlow TTP

- Parse graph implicit in OpenFlow
 - Or at least unspecified...
- Packet formats constrained by table pattern
 - Explicit match on protocol ID with fixed value
 - OpenFlow match field prerequisites
 - Can be further constrained by **flow_paths**





air_meta.yaml

- **parser** provides a parse graph and protocol ID bindings
 - Parse graph is a directed graph connecting **parse_state** values
 - Binds protocol ID field values (**select_value**)
 - Also **value set** can be referenced in parser
 - Enter parser and reach some parse state
 - Accepting and exception states (no exception feature yet)
 - Constrains packet structure
 - Prevents run-time binding of protocol ID values
 - E.g. Controller-defined Ethernet TPID
- Can the same information be derived from table definitions?
- Can the same constraints be provided in table definitions?
 - Depends on run-time flexibility in goto_table?
 - Use built-in table entries? Read-only tables?





From discussion on 6 Feb 15 call :

- Also may need to indicate nesting level of a header
 - Each instance is explicitly defined by a separate **header** definition
 - May need more flexible scheme that allows reuse by not requiring a distinct header name per “level”
 - Three orthogonal concepts: field name, protocol ID value, location in packet
 - How is complexity of required parsing indicated in the IR?
 - How would MPLS entropy label be represented (two RFCs for this)?
- Tradeoff between
 - Requiring a distinct name for each field location in (any) packet
 - Allowing a field name to be (re)used in different contexts

Parse Graph



OF-PI

- OF-PI also allows run-time binding
 - E.g., as supported by POF and some hardware platforms

How can the IR support both
config-time binding and run-time binding?



Mechanisms for composing actions

- Action list (**APPLY_ACTIONS** instruction)
 - Specified order of actions – config-time/run-time
- Action set (**WRITE_ACTIONS** instruction)
 - Canonical order of actions (datapath dependent) – spec-time
- Table match (MAT function)
 - Select highest priority matching table entry – config-time/run-time
- Table selection (**GOTO_TABLE** instruction)
 - Select next table from match context – config-time/run-time
- Control program
 - Imperative program guides MAT invocation – config-time





OpenFlow TTP

- Table defines {<match, action>} pairs
- Constrains table entry types independent of pipeline location
- Pipeline order controlled by **GOTO_TABLE** instructions
 - But we want this to be more flexible for reusability





air_meta.yaml (first version)

- **table** defines only match condition
 - Do we assume multi-field match allows for arbitrary subsets?
 - Match types exact, ternary, valid
 - Separate logic from tables (e.g., all fields referenced by a table should be valid)
 - Or must this be explicit?
 - Two concerns
 - Specifying constraints explicitly
 - Transformation to an efficient implementation
 - If valid matches are A+B and A+C but not B+C, for example
 - Port+EtherType -> VNI (untagged ports)
 - Port+VID -> VNI (tagged ports)
 - Mapping complex function into a single table
 - Field A exact table, what if A is not valid? Various ambiguities. How to address these?
 - How exact must be IR be regarding run-time constraints on the datapath behavior?
 - Could use table matching on Port+VID+EtherType, but
 - Would this allow match on EtherType+VID? (invalid case)
 - Does preventing this require two tables?



air_meta.yaml (current version)

- **table** defines match condition and valid actions
 - Does this provide sufficient constraint information?
 - Still relies on **control_flow** to provide additional constraints?



air_meta.yaml (first version)

- **control_flow** defines match+action pair in pipeline context
 - Can a table be consulted multiple times in the pipeline?
 - Nothing prevents this currently; could add a constraint to prevent this
 - How does this constrain table entries?
 - E.g., what if a table entry carries a different action at run-time?
 - Run-time error checking (control context dependency)
 - Control flow decisions may depend on match context
 - Requires additional metadata

