



# Protocol Independent Forwarding

ONF PIF Project

2015-02-13

Underlined text records discussion

# Agenda




- Overview (30 min)
  - Motivation and terminology
  - Projects and responsibilities
- Tour of AIR-IRI (30 min)
  - Sample datapath program
  - Overview of interpreter + infrastructure
- Work items and infrastructure (40 min)
  - Summary of work so far
  - Future work areas - IR development (interpreter+tools), IR samples, runtime interacting with IR
  - Software infrastructure and licensing
- Next steps - how to participate (20 min)

# Why OpenFlow Next Generation / PIF etc?





- Vision - refactor architecture / design (*technology push*)
  - Support new approaches, e.g. more flexibly configured datapath => *Protocol Independent Forwarding*
  - Support new operations, e.g. stateful, delegated to switch, nested encapsulation
  - Optimize capacity, throughput - e.g. simpler, more parallel
- Optimize processes / lifecycles (*facilitate ecosystem*)
  - Standards creation process (help ONF working group participants)
  - Development and deployment processes (help switch, controller, app vendors and users / operators)
    - Avoid cross product of controllers / apps and switch types
- Address specific requirements (*market pull*)
  - Support new use cases / market segments e.g. L4-L7 services / NFV
  - Refinement of capabilities / improved performance etc. for existing use cases

# Interpretations of Protocol Independence



- Restructuring the existing OpenFlow *specification* - result:
  - Modular specification
  - Core or base does not refer to protocols
  - Each protocol / layer (e.g. Ethernet or IP) documented in its own add on
  - Result can express the same semantics as e.g. OpenFlow 1.3 - no impact on switches / controllers
  - Easier to add support for protocols in future - just write an add on module
- Introducing support for “*user defined*” *protocols* - result:
  - Match fields not limited to existing set of 40-ish OXMs
  - “Users” (vendors / operators...) describe new field as length + offset from already defined field etc.
  - Can be accomplished by extending existing OpenFlow 1.x specification
- More ambitious Protocol Independent Forwarding project - result:
  - More than just ability to define new protocol fields
  - Enables defining arrangement of OpenFlow pipeline in a more flexible way
  - “Datapath program” (in effect forwarding model) describes pipeline arrangement (matching tables, actions, QoS TM elements etc.) and behavior
  - Toolchains (compilers etc) enable configuring switches to “run” these “programs”

# Elements in PIF World - New + Improved

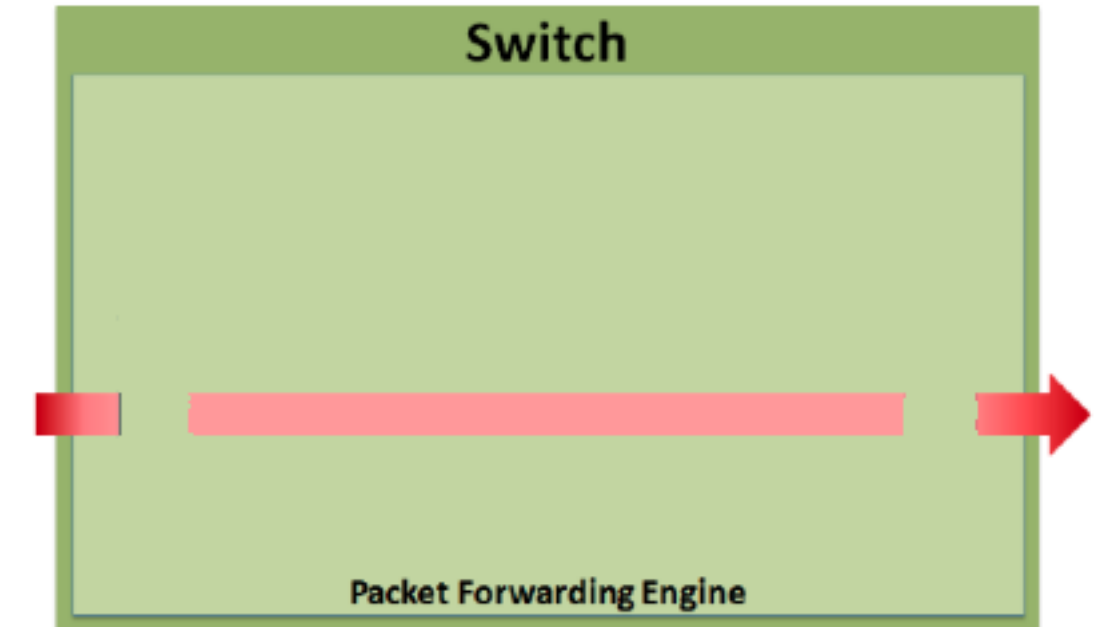



- Predefined Protocol Forwarding
  - OpenFlow specification defines protocols / fields, match/action behavior, overall control flow (tables can influence)
  - Set of supported protocols fixed by implementation

# Elements in PIF World - New + Improved

- Config time



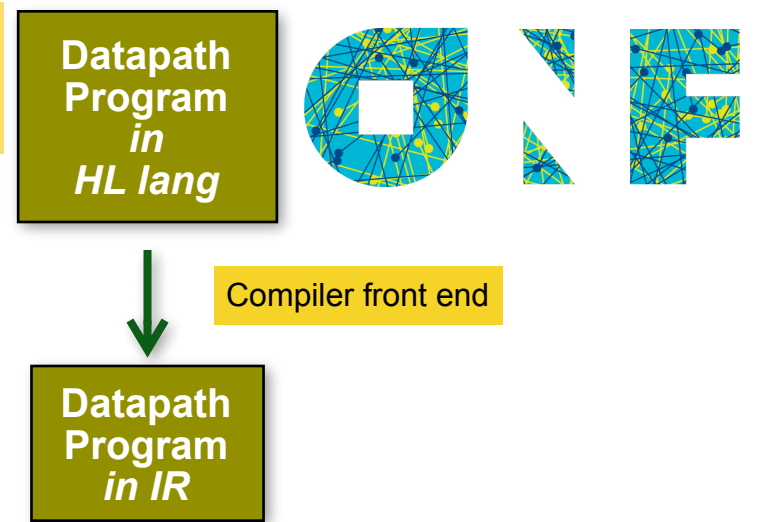
Datapath  
Program  
in  
HL lang



- Predefined Protocol Forwarding
  - OpenFlow specification defines protocols / fields, match/action behavior, overall control flow (tables can influence)
  - Set of supported protocols fixed by implementation

# Elements in PIF World - New + Improved

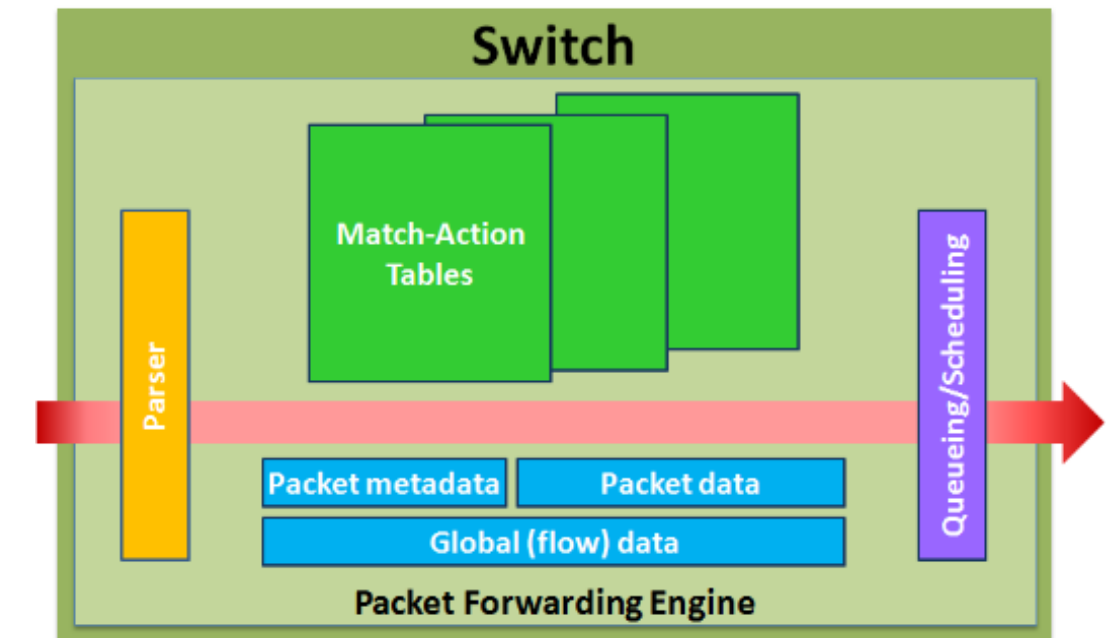

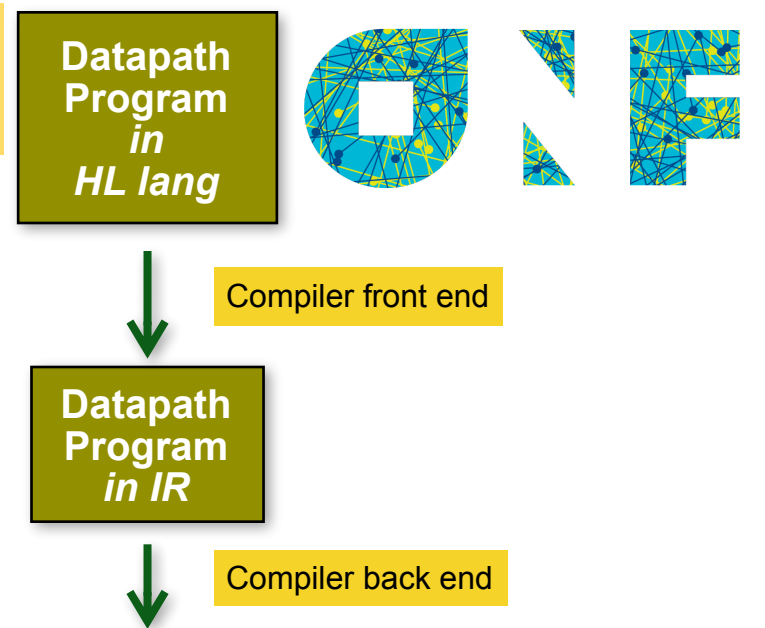
- Config time



- Predefined Protocol Forwarding
  - OpenFlow specification defines protocols / fields, match/action behavior, overall control flow (tables can influence)
  - Set of supported protocols fixed by implementation

# Elements in PIF World - New + Improved

- Config time



- Predefined Protocol Forwarding
  - OpenFlow specification defines protocols / fields, match/action behavior, overall control flow (tables can influence)
  - Set of supported protocols fixed by implementation



# Elements in PIF World - New + Improved

- Config time

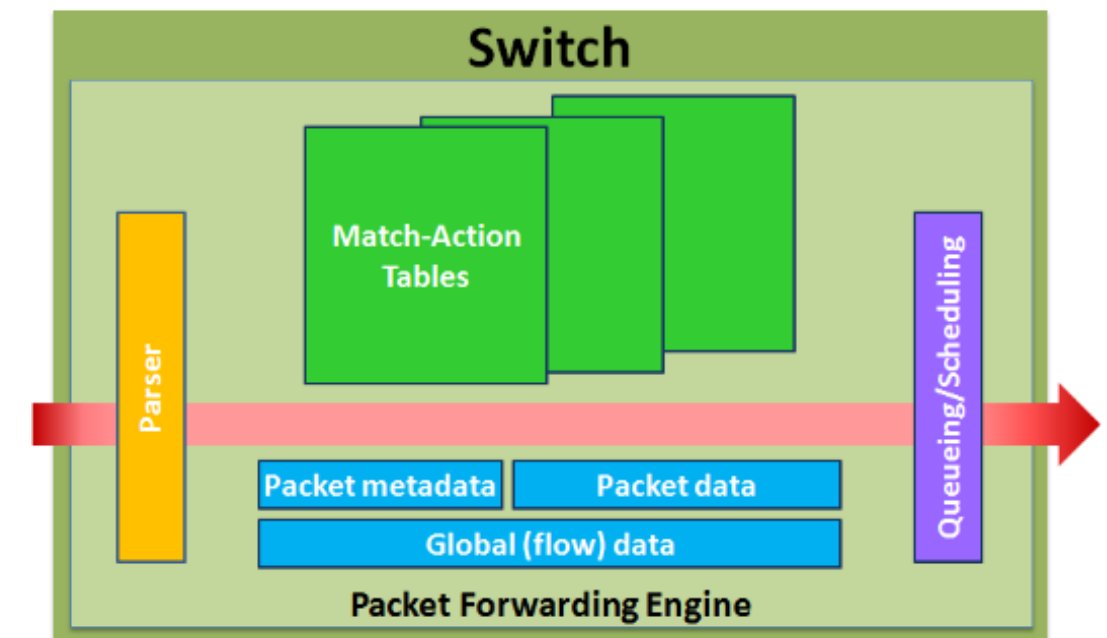
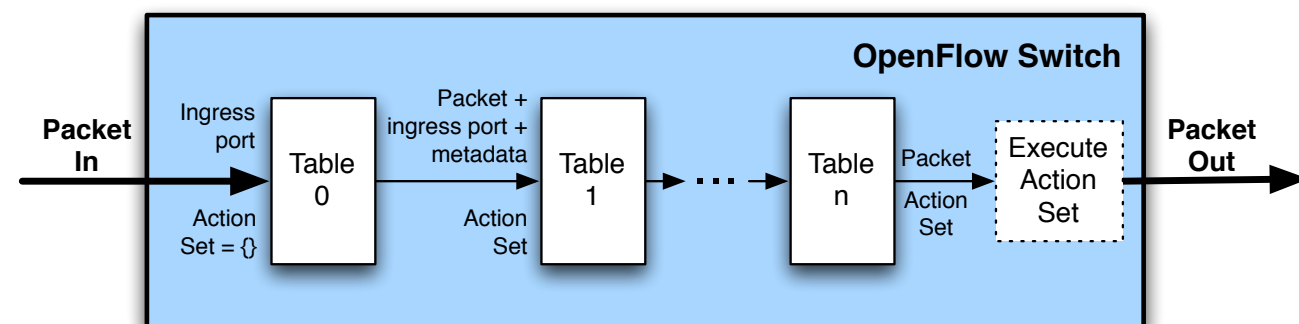
Datapath Program  
in  
HL lang



Compiler front end

Datapath Program  
in IR

Compiler back end



- Predefined Protocol Forwarding
  - OpenFlow specification defines protocols / fields, match/action behavior, overall control flow (tables can influence)
  - Set of supported protocols fixed by implementation

- Protocol Independent Forwarding (PIF) Configured (“Programmed”) Datapath
  - Programs in language(s) describe datapath
    - Parse tree => protocol independent
    - Match/action tables (control flow arranges table sequence)
    - Packet metadata, per table or global state
    - QoS

# Elements in PIF World - New + Improved

- Config time

Datapath Program  
in  
HL lang

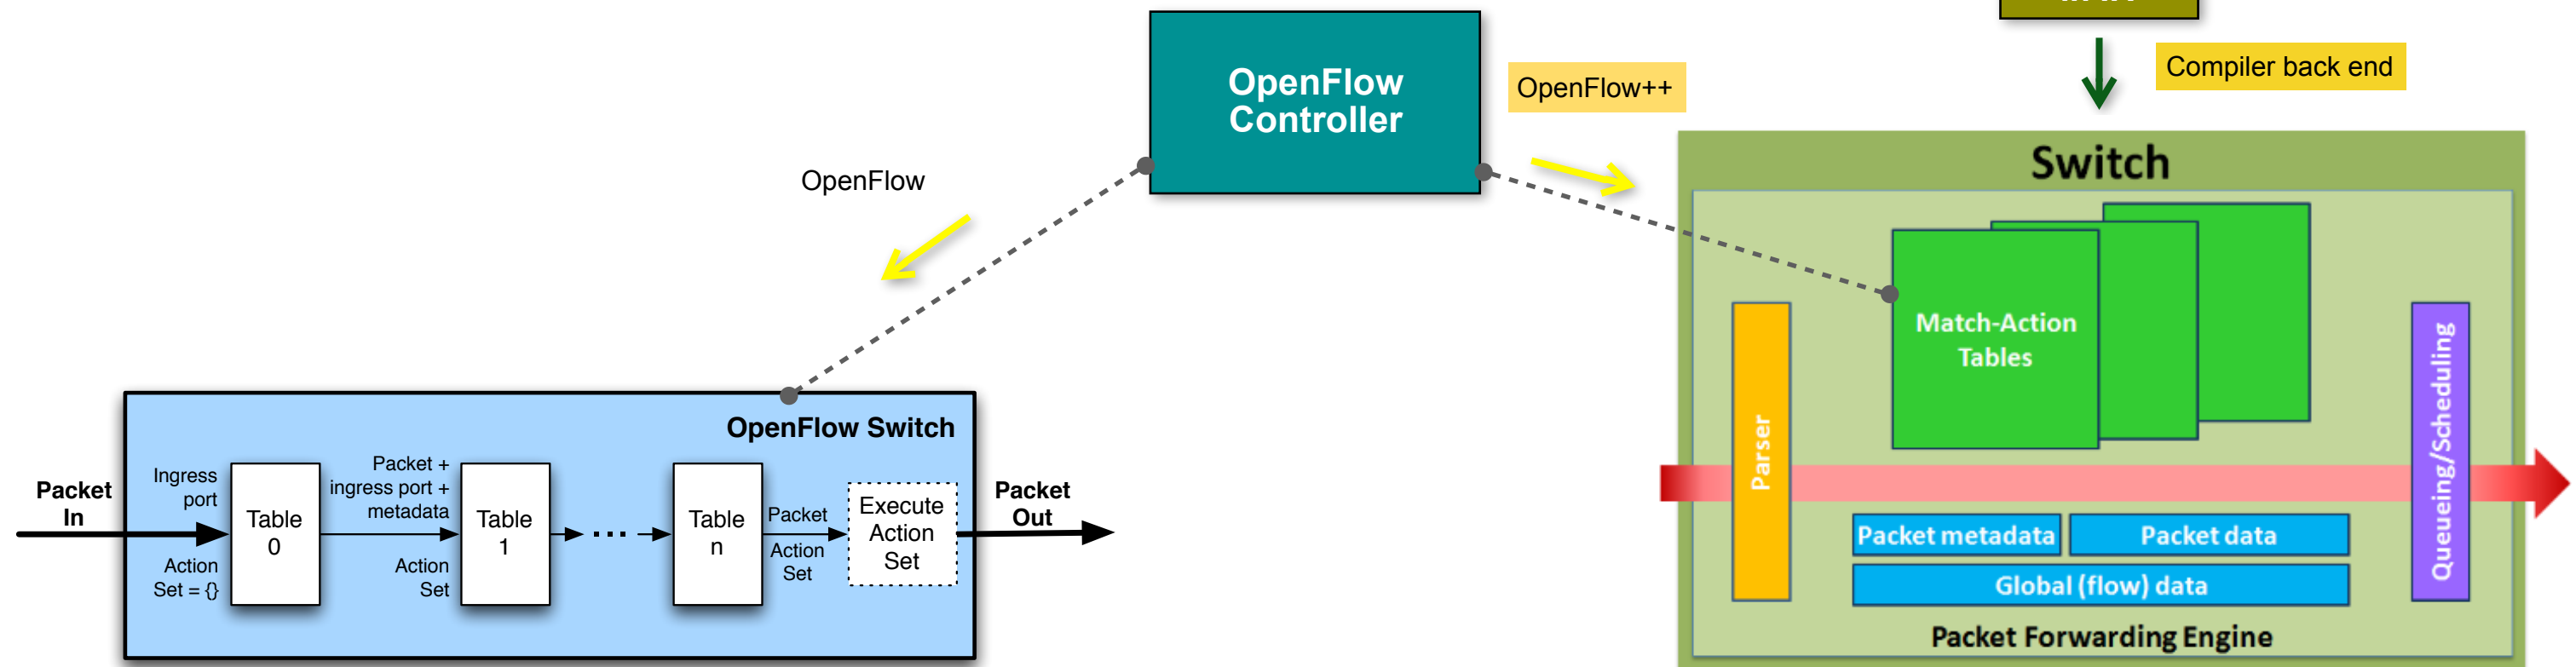


Compiler front end

Datapath Program  
in IR

Compiler back end


- Either way - need “southbound” interface for run-time interaction with switch...
- Populate tables, receive statistics + events



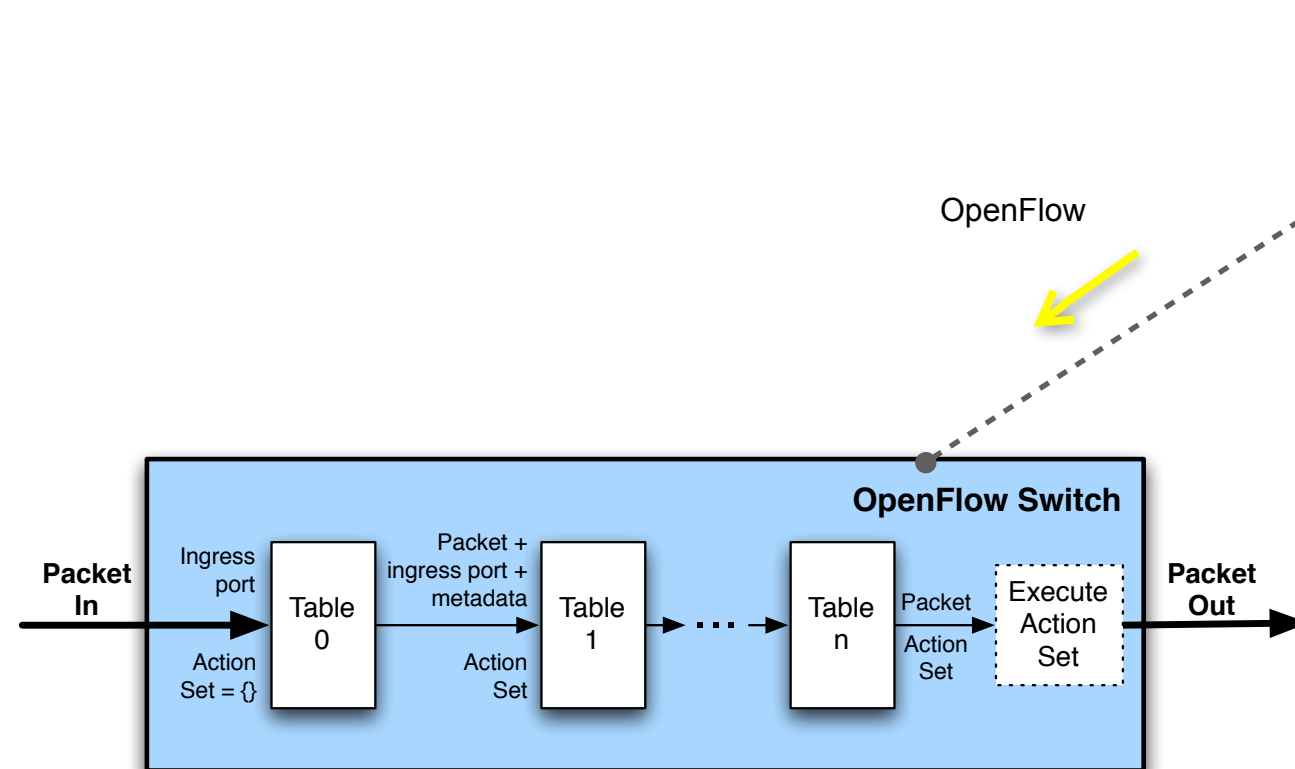
- Predefined Protocol Forwarding
  - OpenFlow specification defines protocols / fields, match/action behavior, overall control flow (tables can influence)
  - Set of supported protocols fixed by implementation

- Protocol Independent Forwarding (PIF) Configured (“Programmed”) Datapath
  - Programs in language(s) describe datapath
    - Parse tree => protocol independent
    - Match/action tables (control flow arranges table sequence)
    - Packet metadata, per table or global state
    - QoS

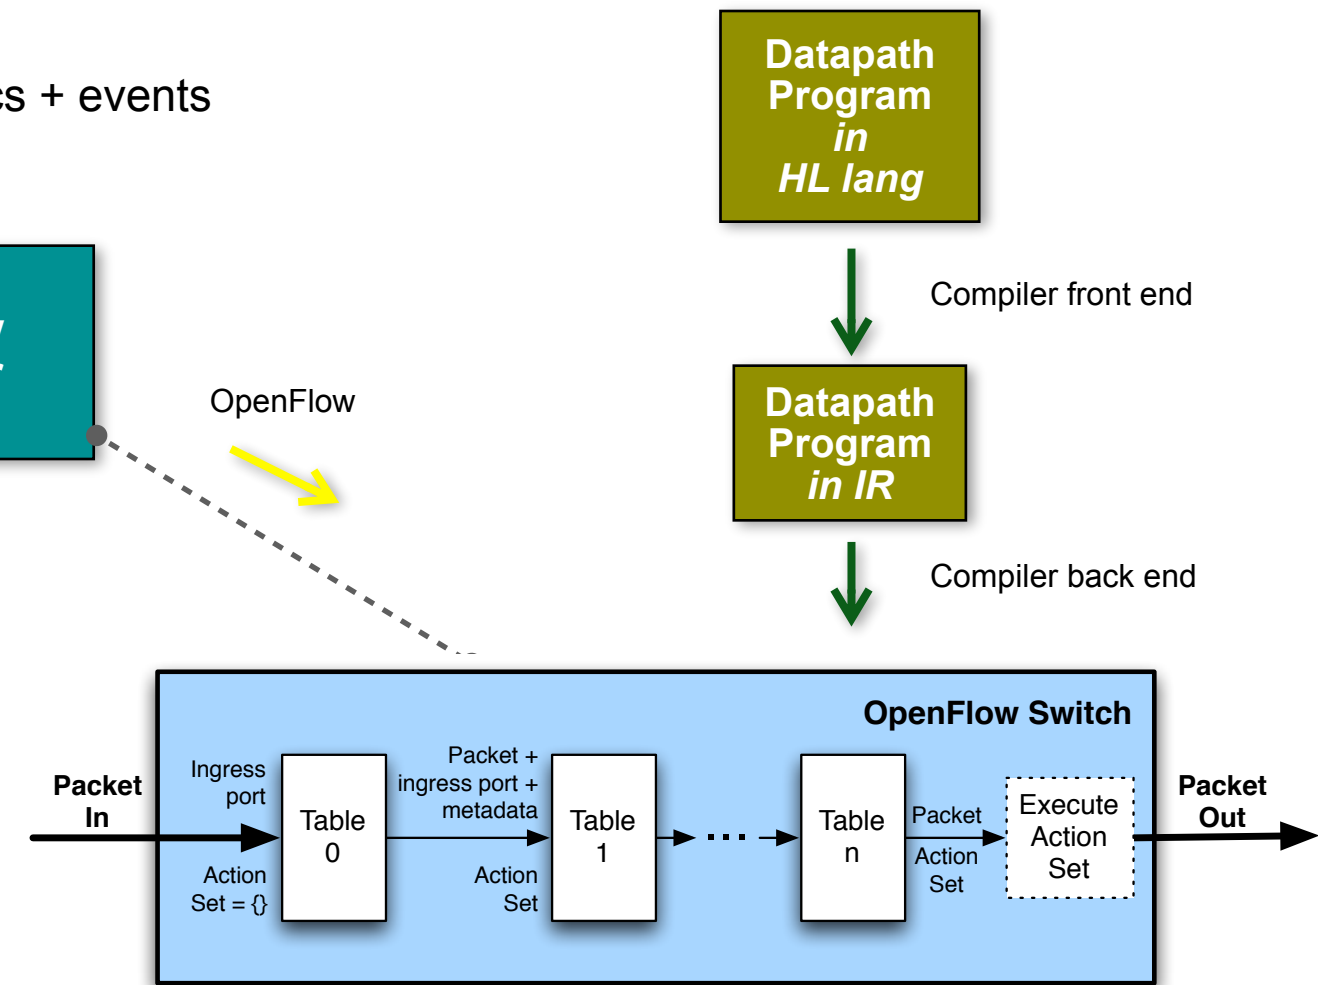
# Emulating OpenFlow 1.x



- Southbound
  - Populate tables, receive statistics + events

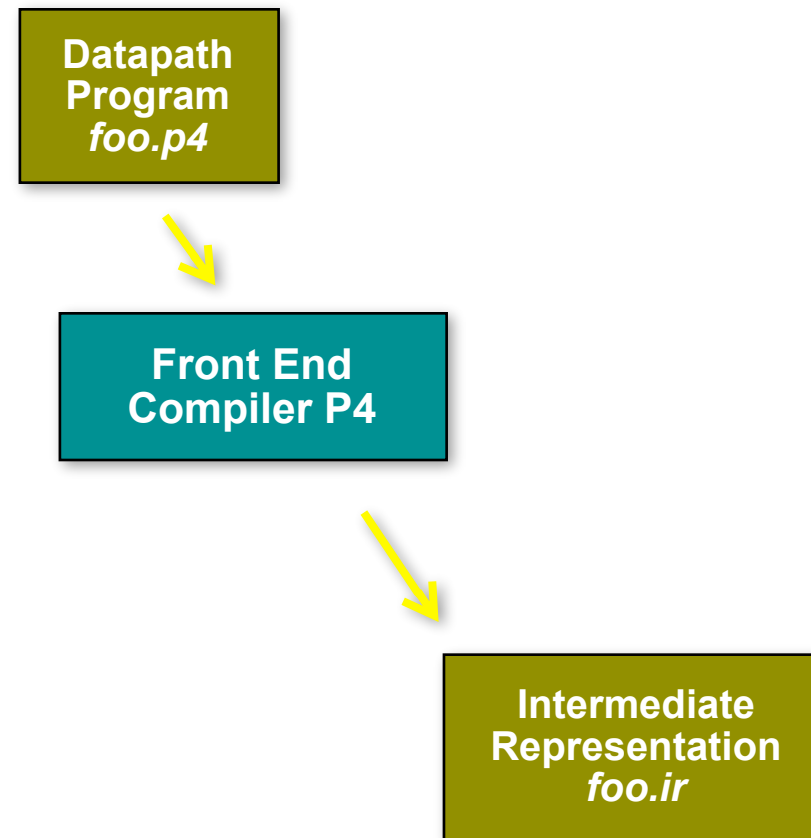


- Predefined Protocol Forwarding
  - Implements OpenFlow 1.x spec manually

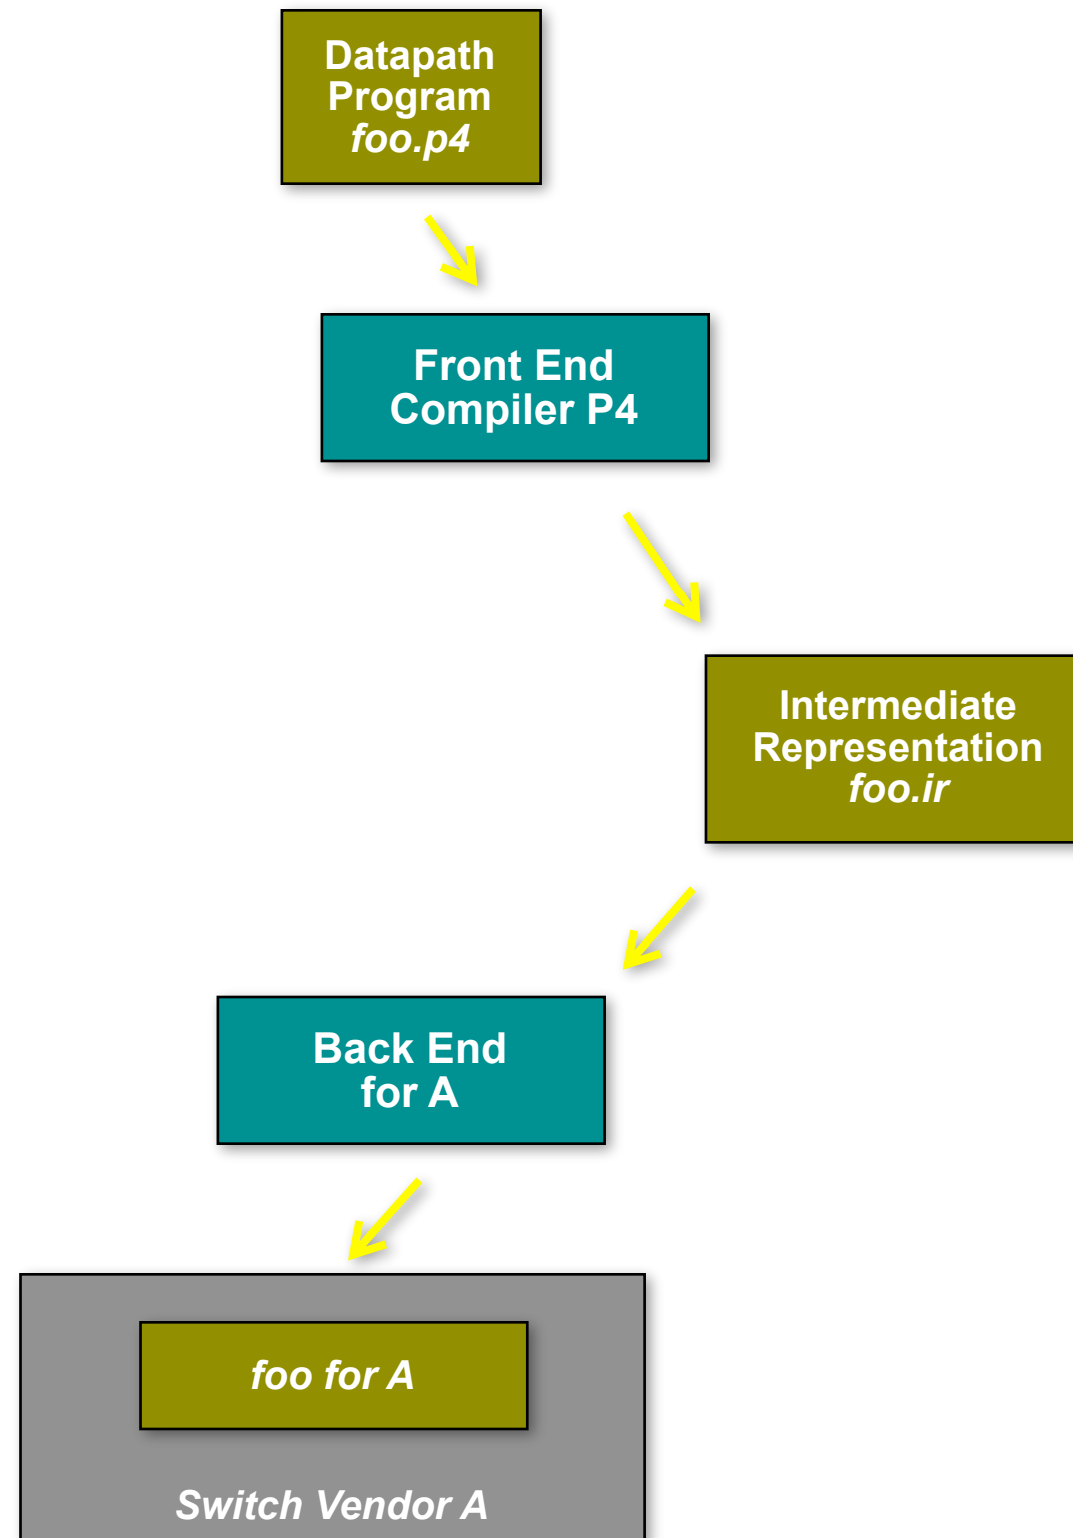


- Configured (“Programmed”) Datapath
  - Program implements OpenFlow 1.x spec

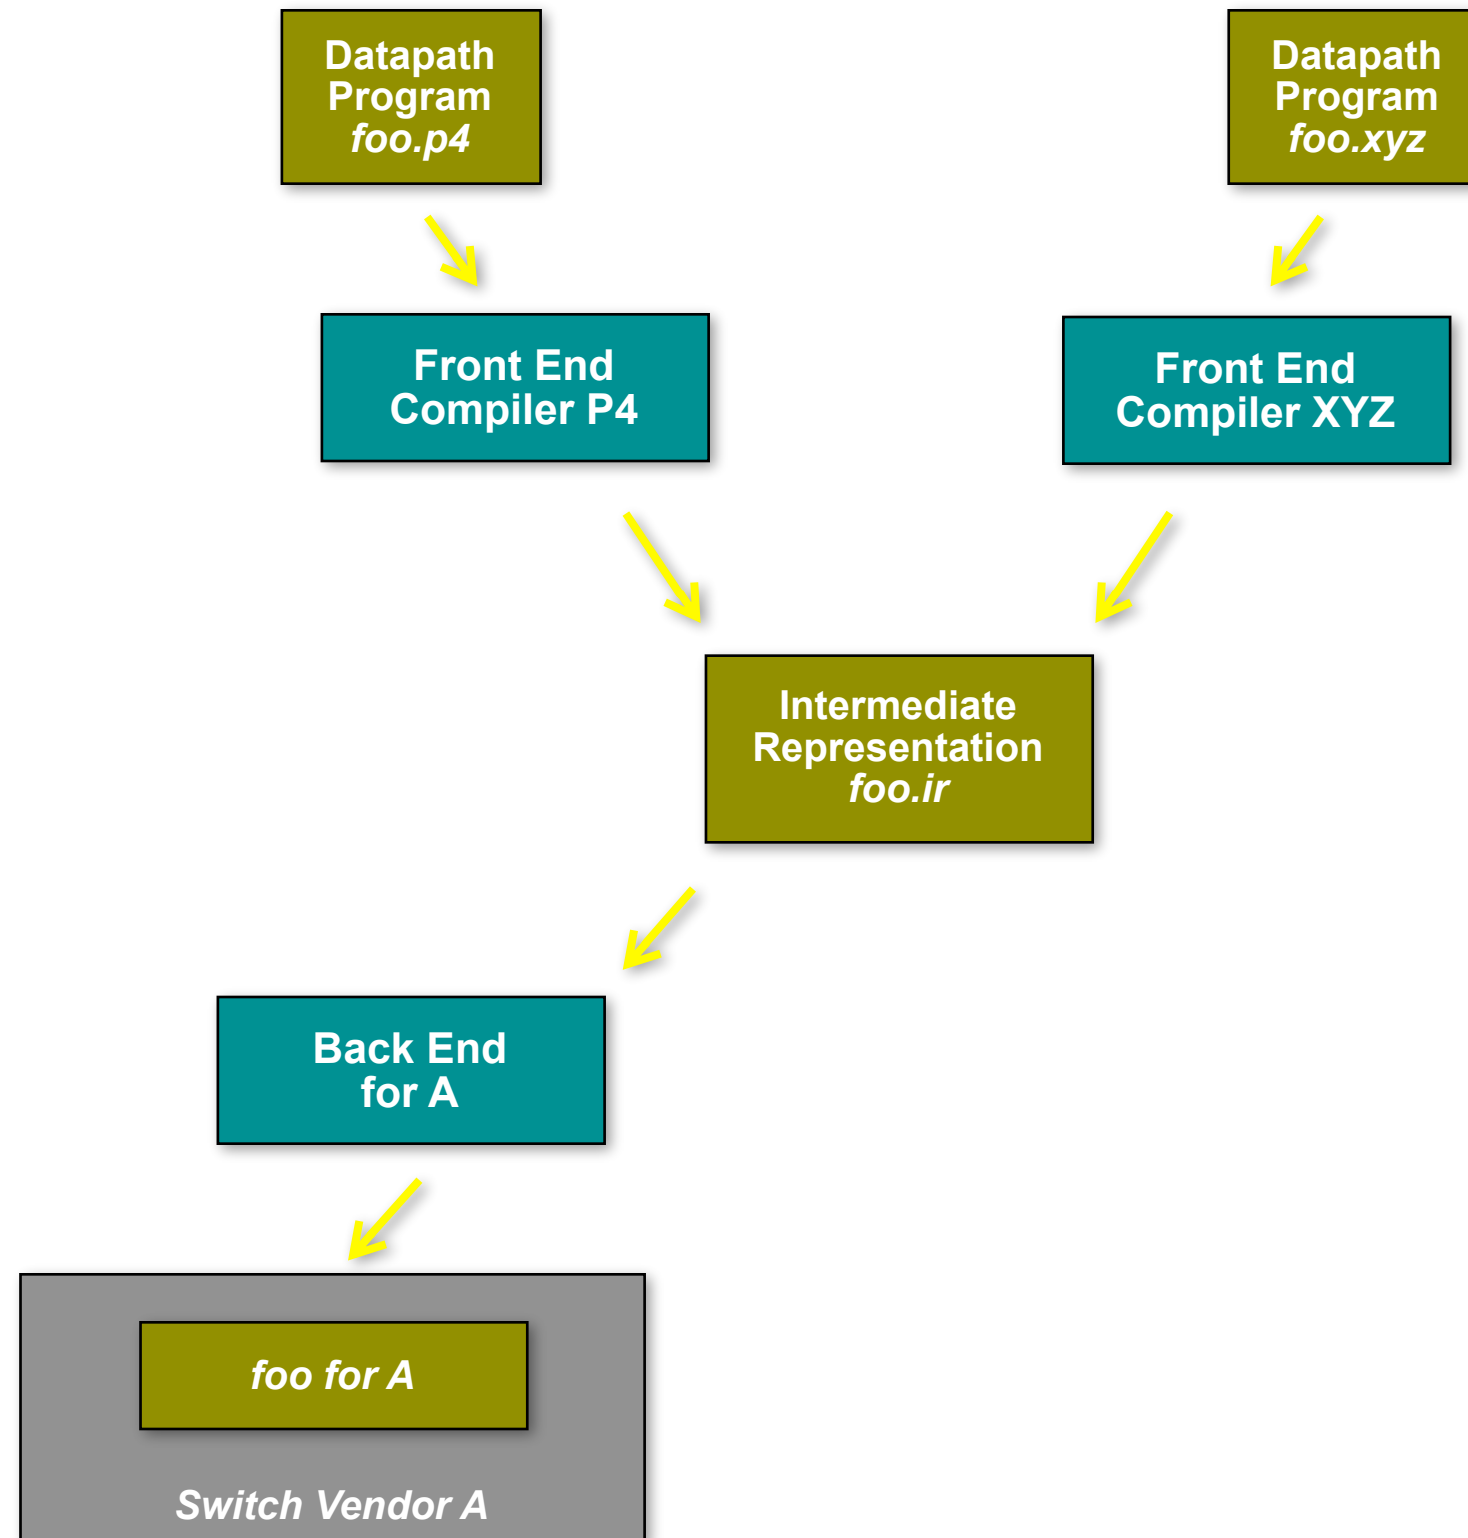
# Elements and Responsibilities




# Elements and Responsibilities



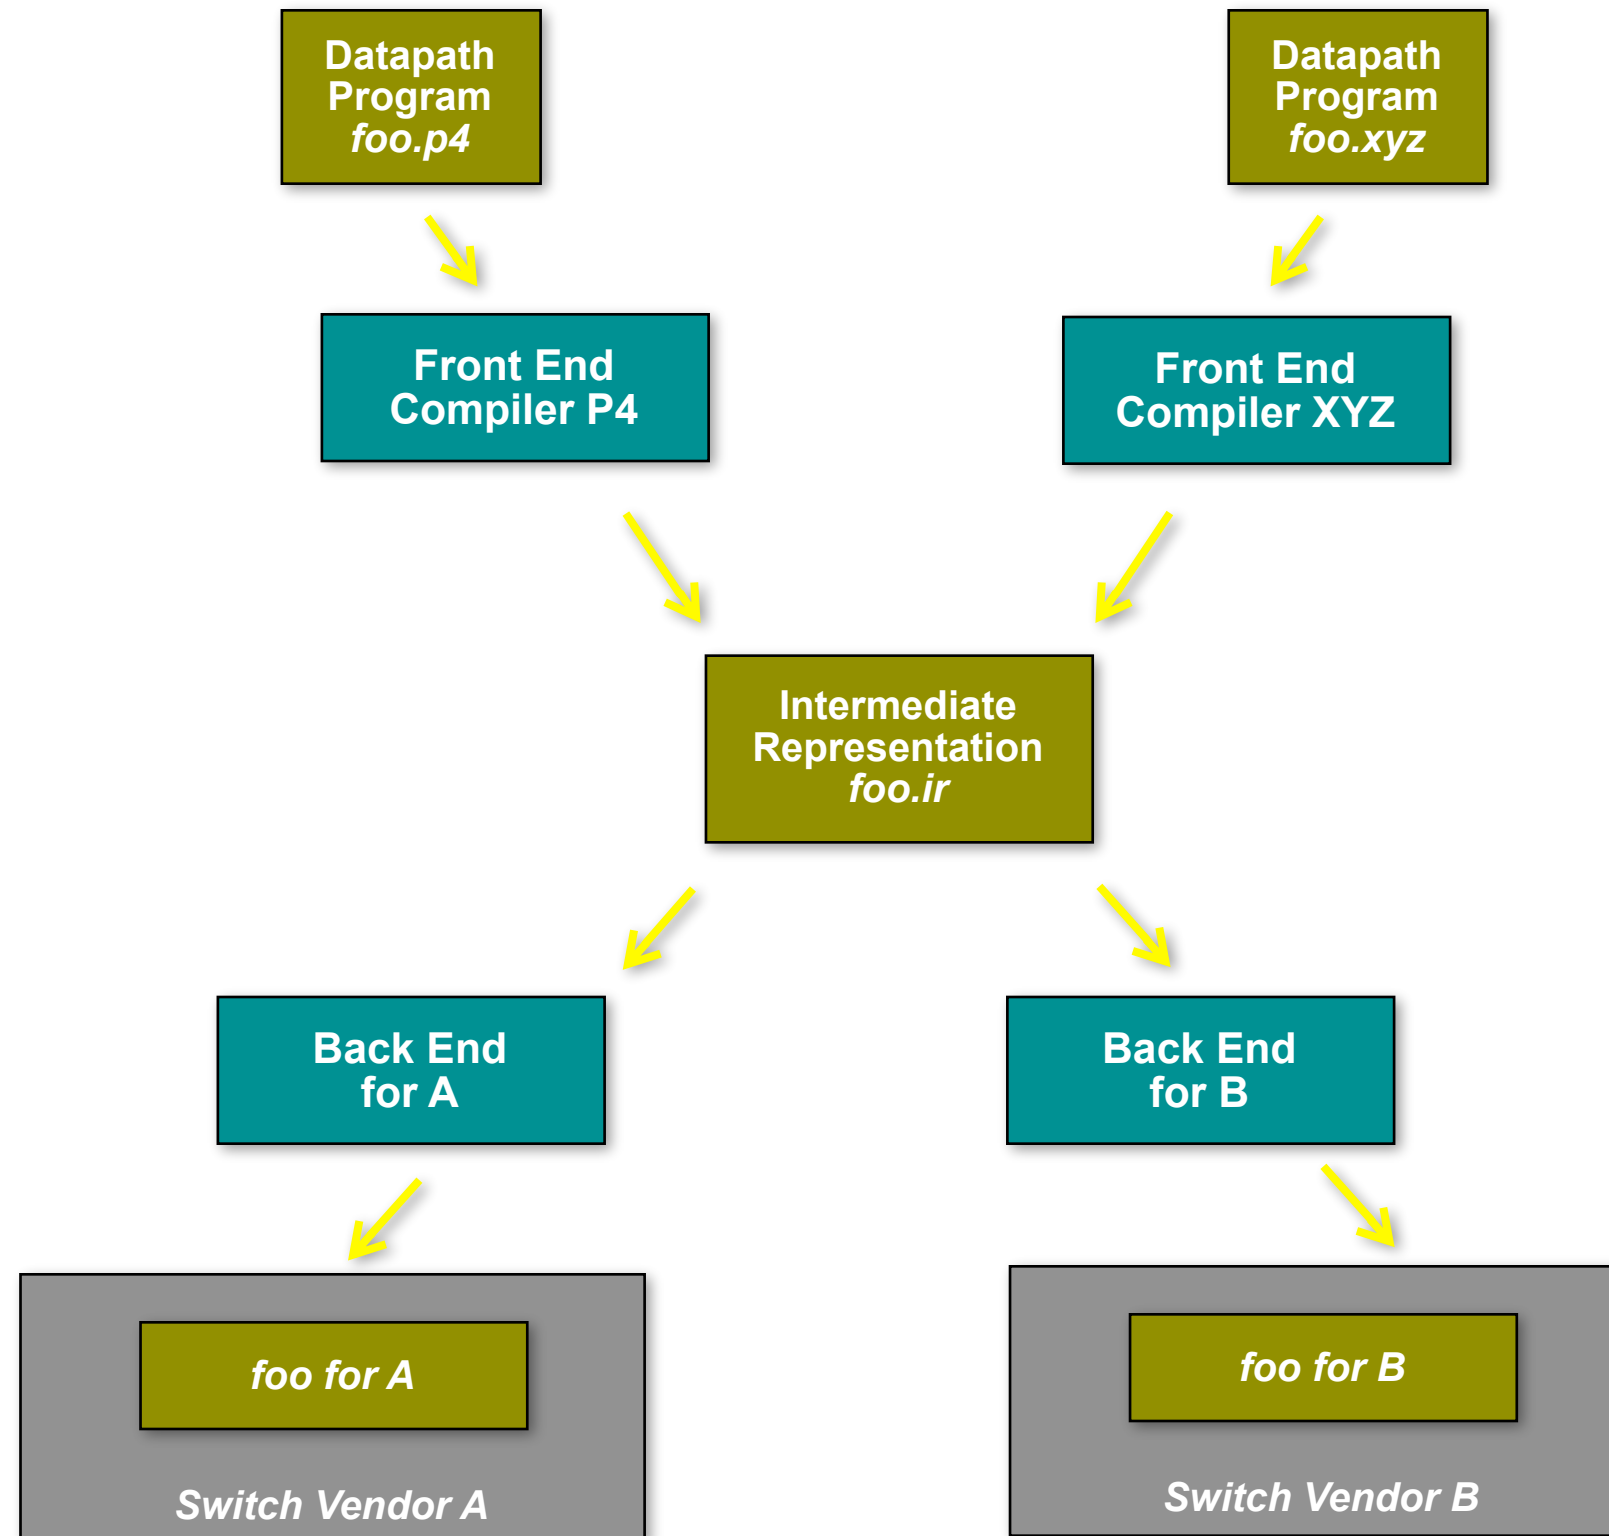
# Elements and Responsibilities



# Elements and Responsibilities



# Elements and Responsibilities



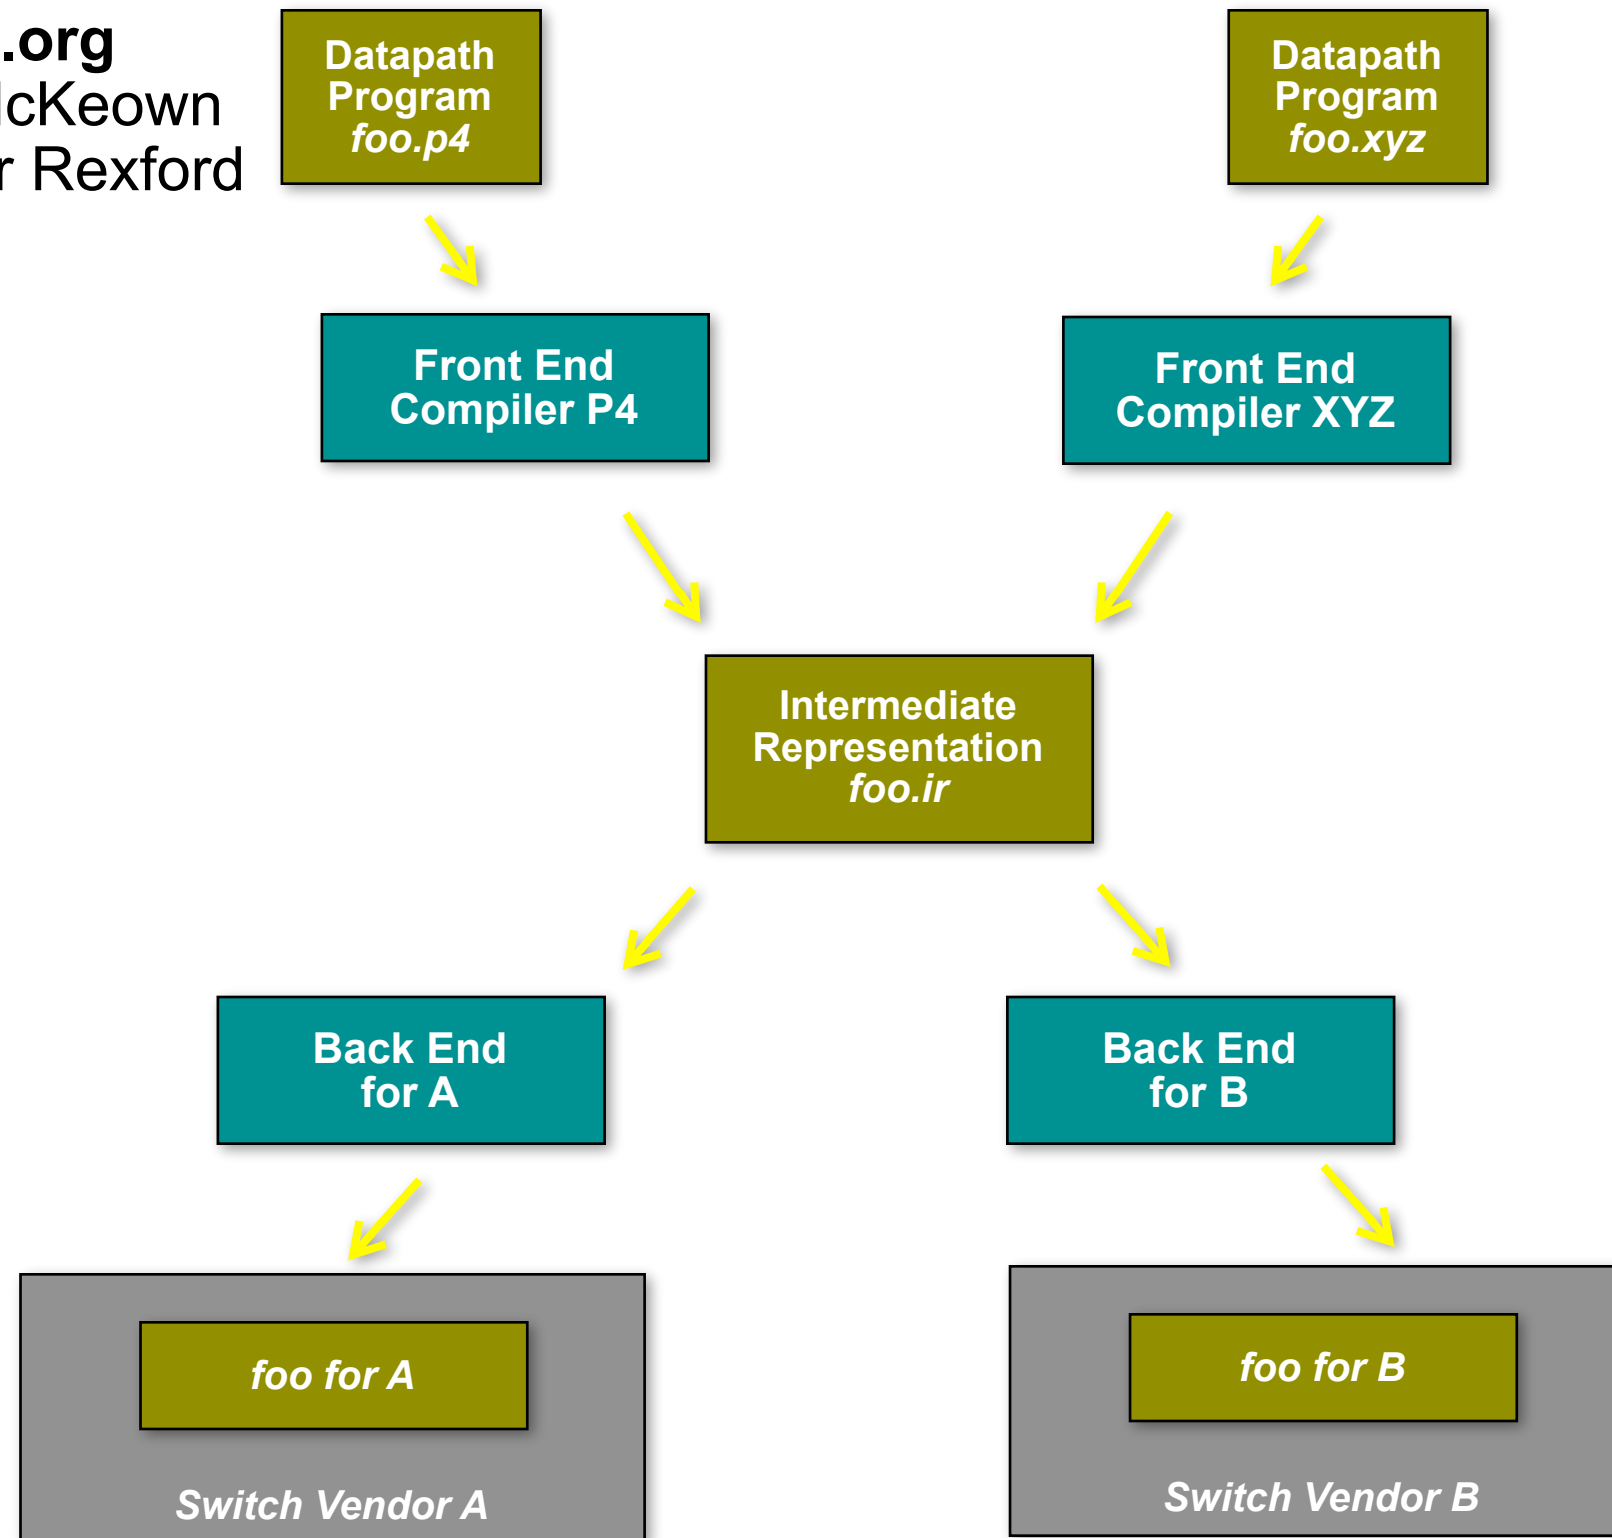
- Target platforms: various pipeline types, FPGAs, processors - CPU/NPU



# Elements and Responsibilities




**P4.org**  
Nick McKeown  
Jennifer Rexford



- Target platforms: various pipeline types, FPGAs, processors - CPU/NPU

# Elements and Responsibilities



**P4.org**  
Nick McKeown  
Jennifer Rexford

Datapath  
Program  
*foo.p4*

Datapath  
Program  
*foo.xyz*

Front End  
Compiler P4

Front End  
Compiler XYZ


Intermediate  
Representation  
*foo.ir*

**Vendor A**

**Vendor B**

Back End  
for A

Back End  
for B



- Target platforms: various pipeline types, FPGAs, processors - CPU/NPU

# Elements and Responsibilities



**P4.org**  
Nick McKeown  
Jennifer Rexford

Datapath Program  
*foo.p4*

Datapath Program  
*foo.xyz*

Front End  
Compiler P4

Front End  
Compiler XYZ

Intermediate  
Representation  
*foo.ir*

**ONF PIF**  
Open Source  
Project

Create definition of IR by  
evolving implementation

**Vendor A**

**Vendor B**

Back End  
for A

Back End  
for B

- Target platforms: various pipeline types, FPGAs, processors - CPU/NPU

*foo for A*  
  
*Switch Vendor A*

*foo for B*  
  
*Switch Vendor B*

# Elements and Responsibilities



**P4.org**  
Nick McKeown  
Jennifer Rexford

Datapath Program  
*foo.p4*

Datapath Program  
*foo.xyz*

Front End  
Compiler P4

Front End  
Compiler XYZ

Intermediate  
Representation  
*foo.ir*

**ONF PIF**  
Open Source  
Project

Create definition of IR by  
evolving implementation

**Vendor A**

**Vendor B**



Back End  
for A

Back End  
for B

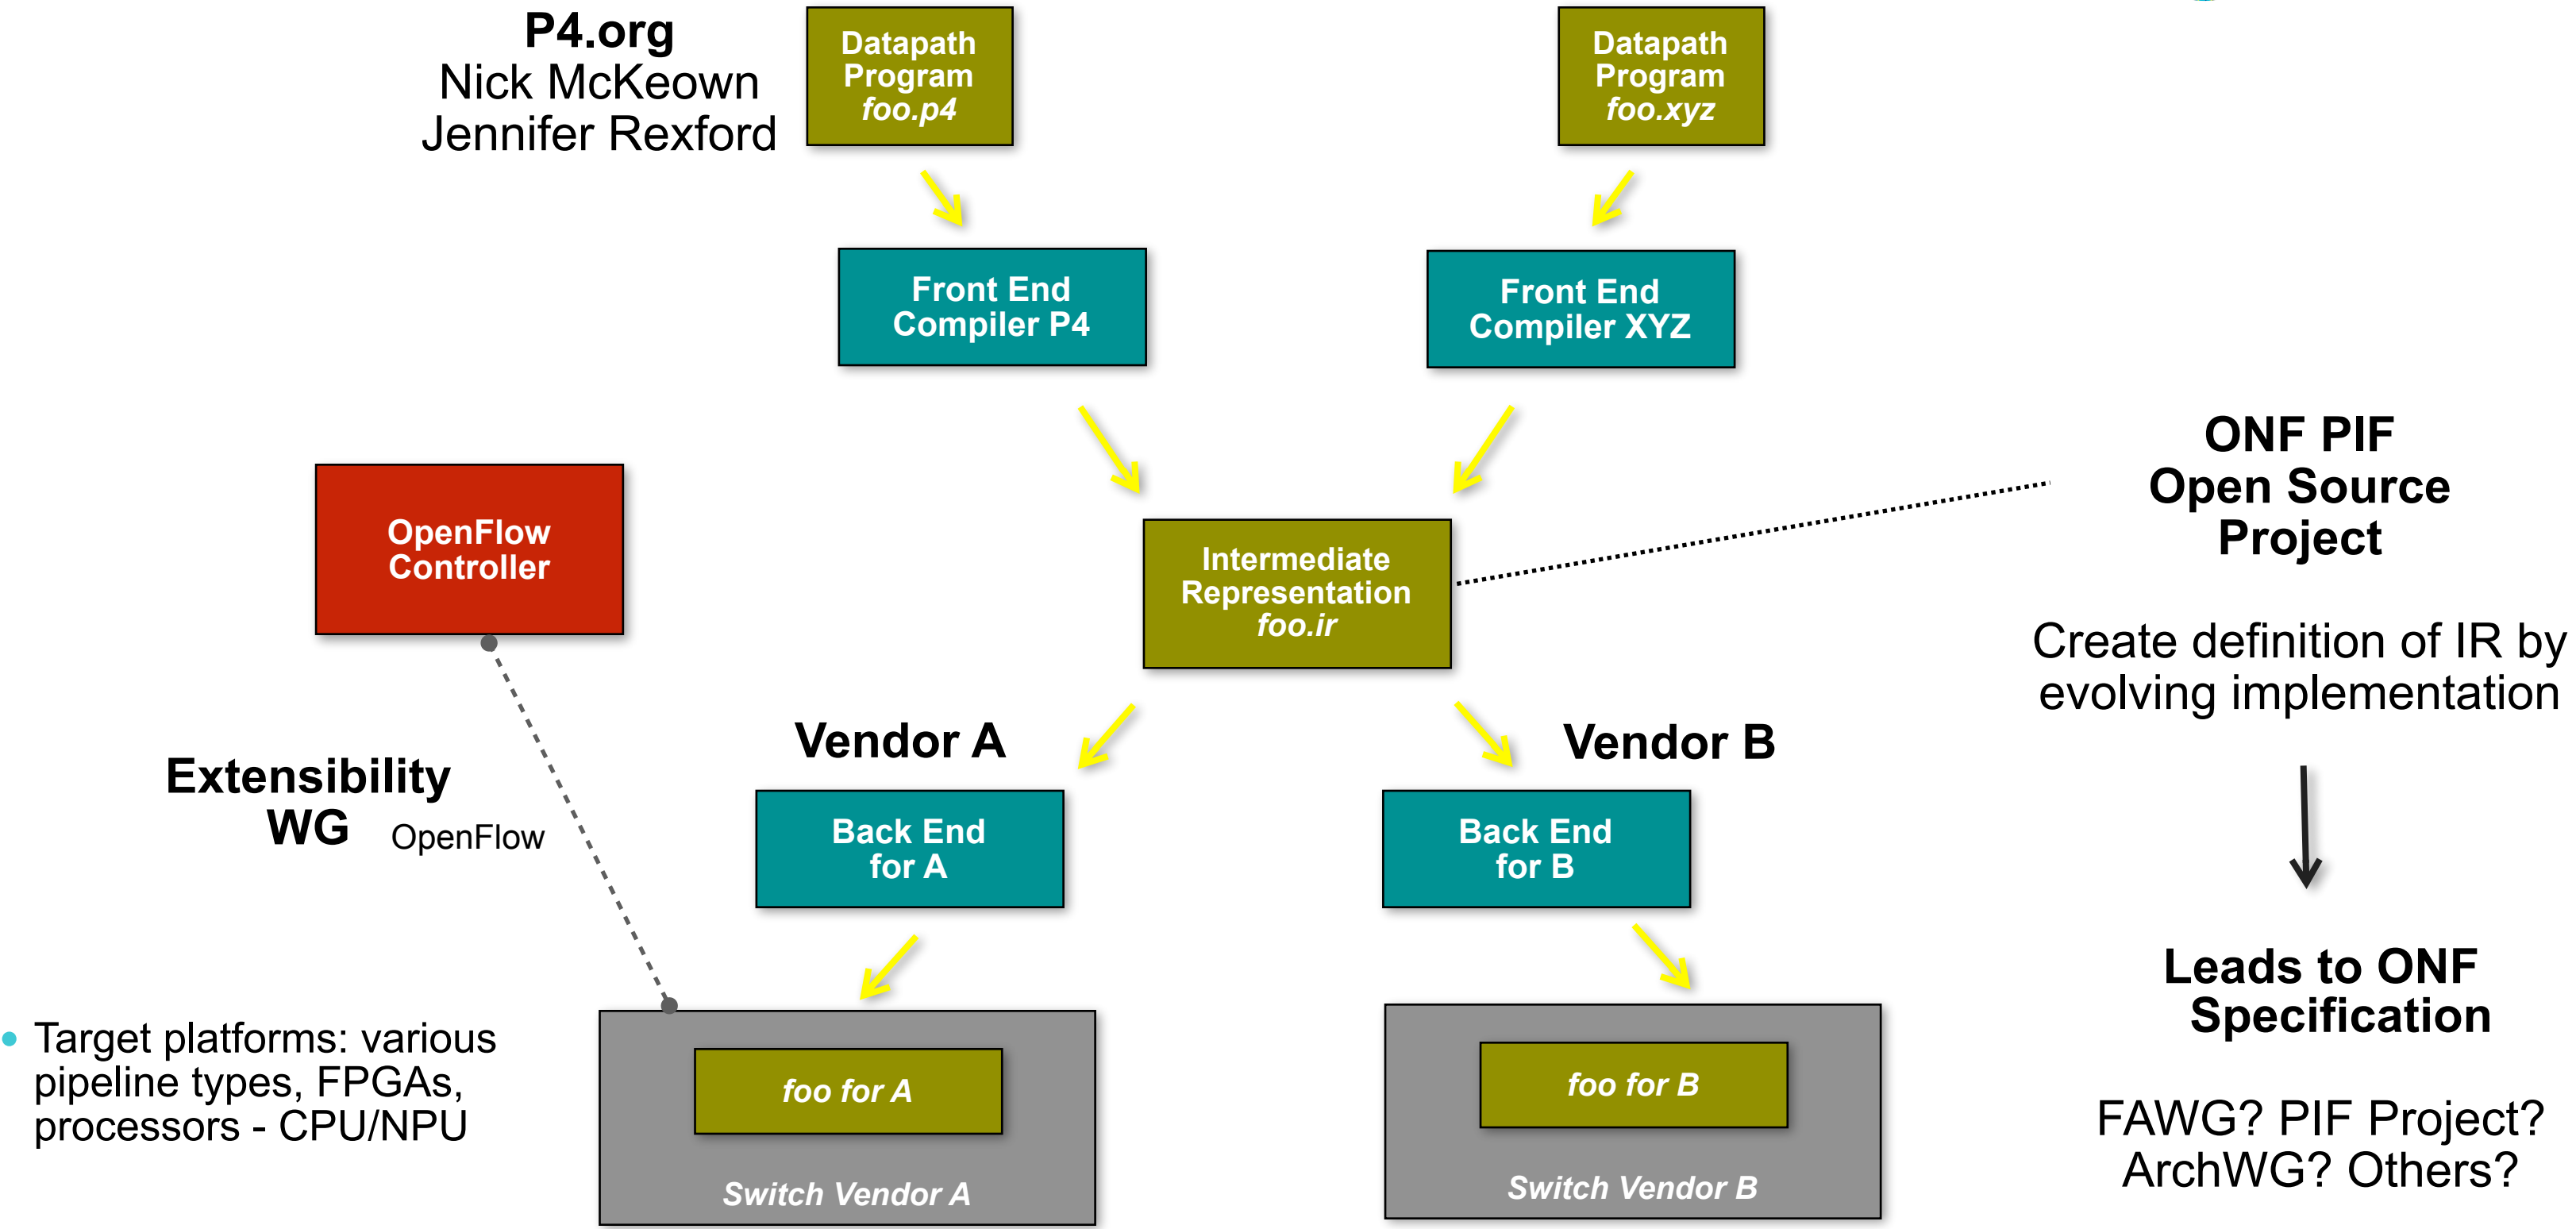
**Leads to ONF  
Specification**

FAWG? PIF Project?  
ArchWG? Others?

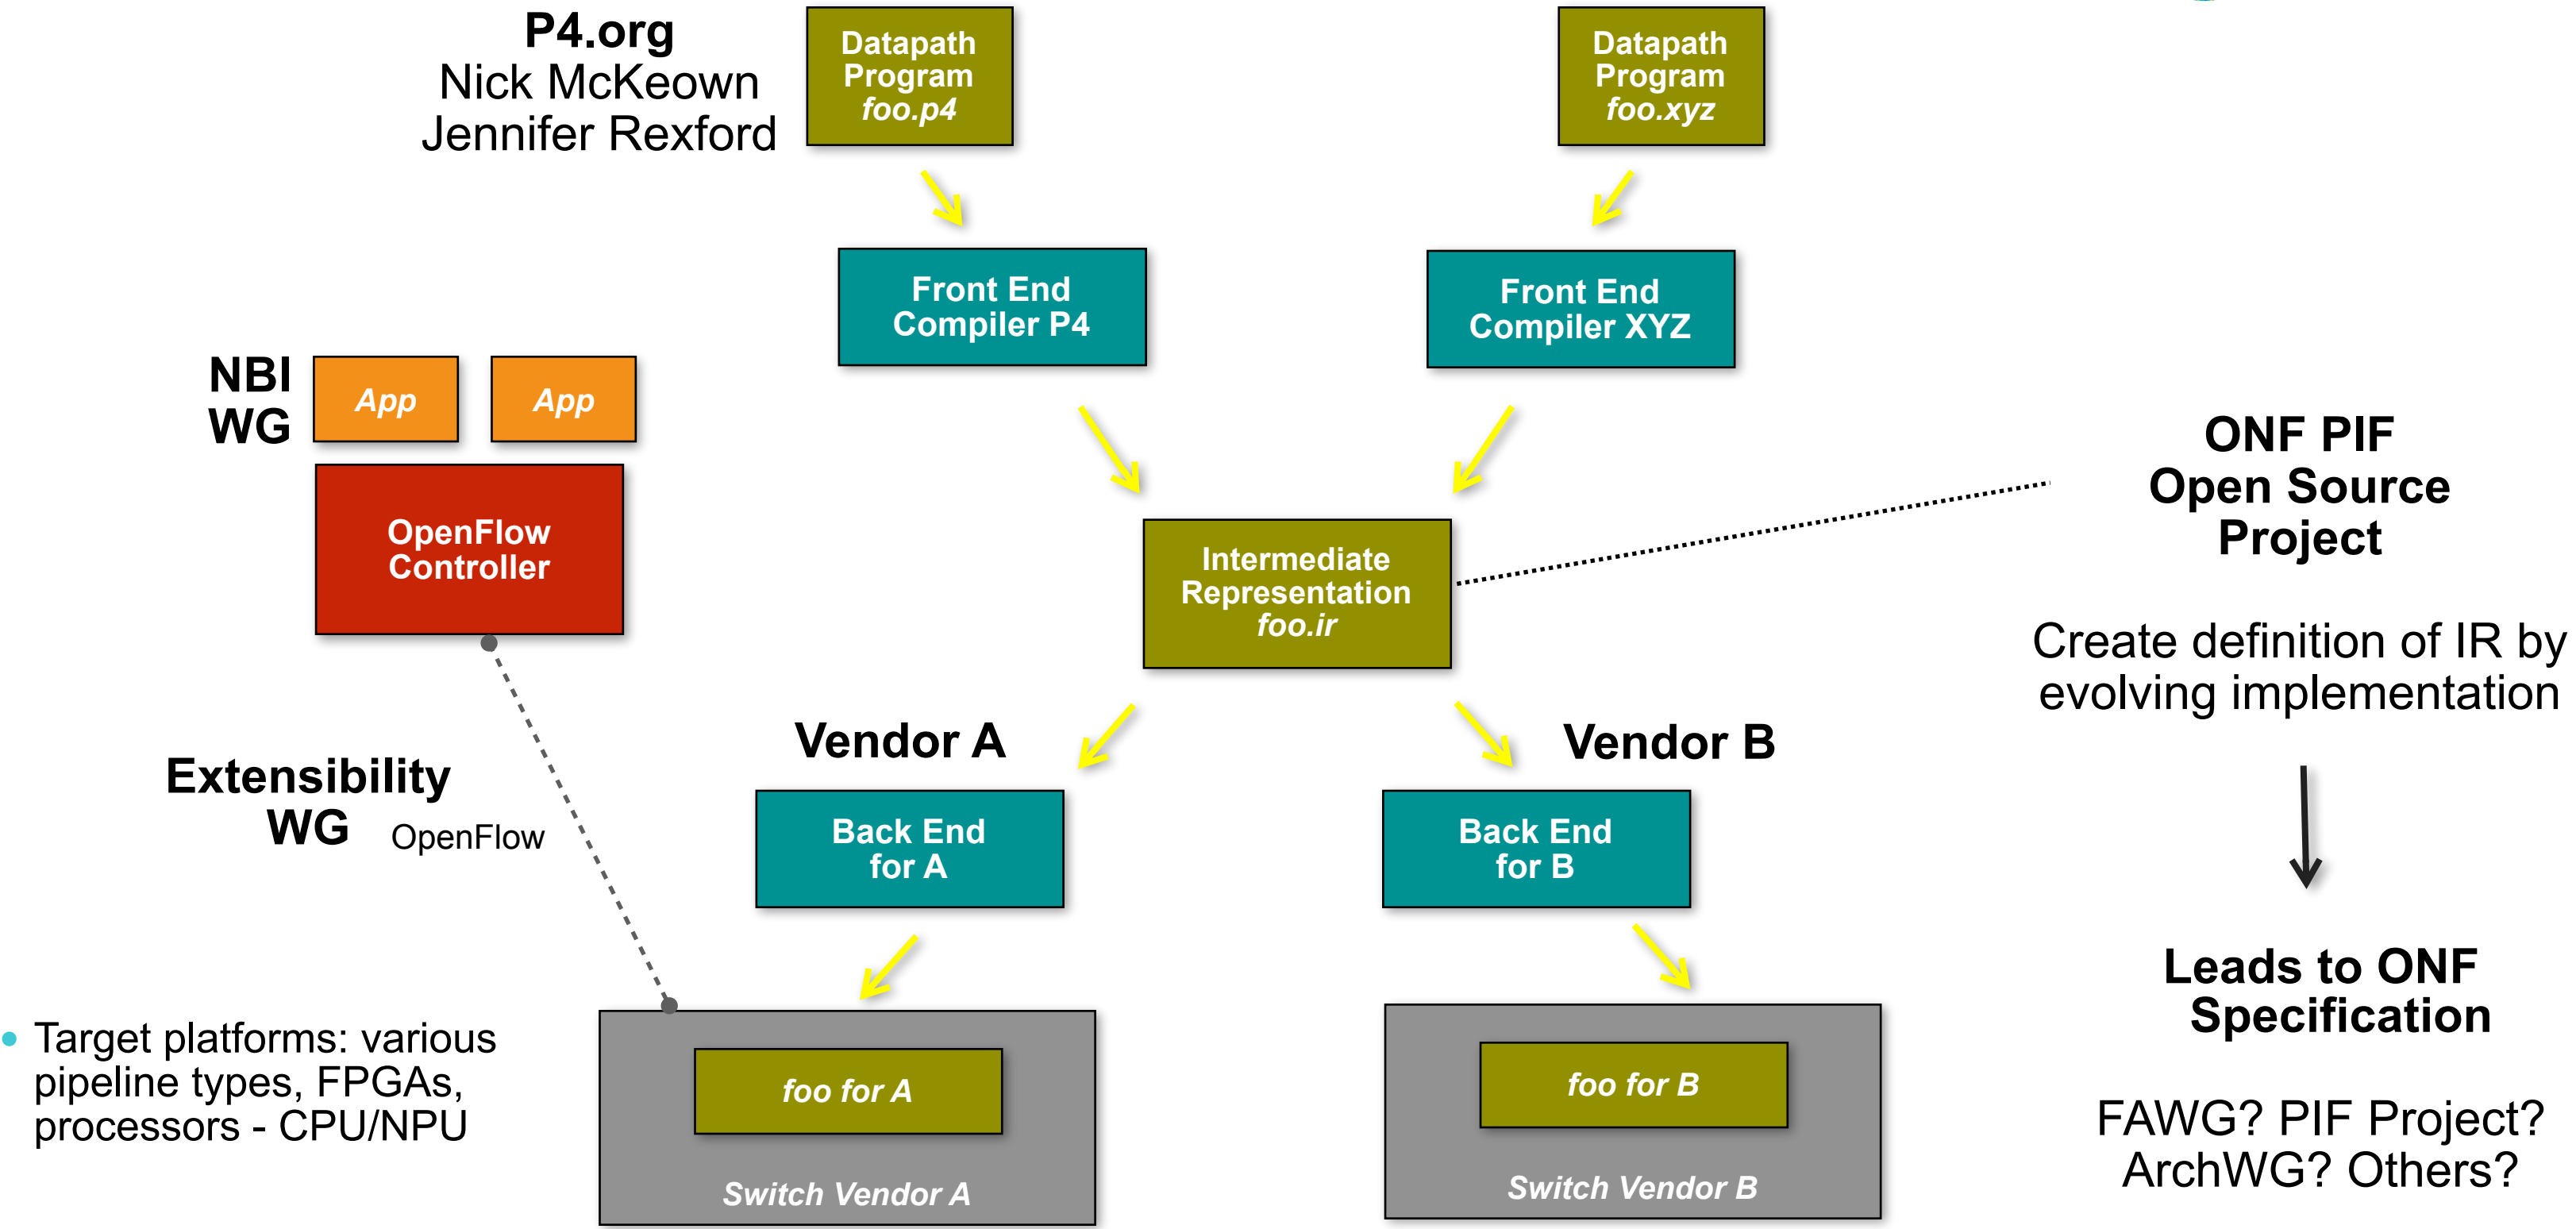
- Target platforms: various pipeline types, FPGAs, processors - CPU/NPU



# Elements and Responsibilities



# Elements and Responsibilities



- Target platforms: various pipeline types, FPGAs, processors - CPU/NPU

# Elements and Responsibilities



## FAWG

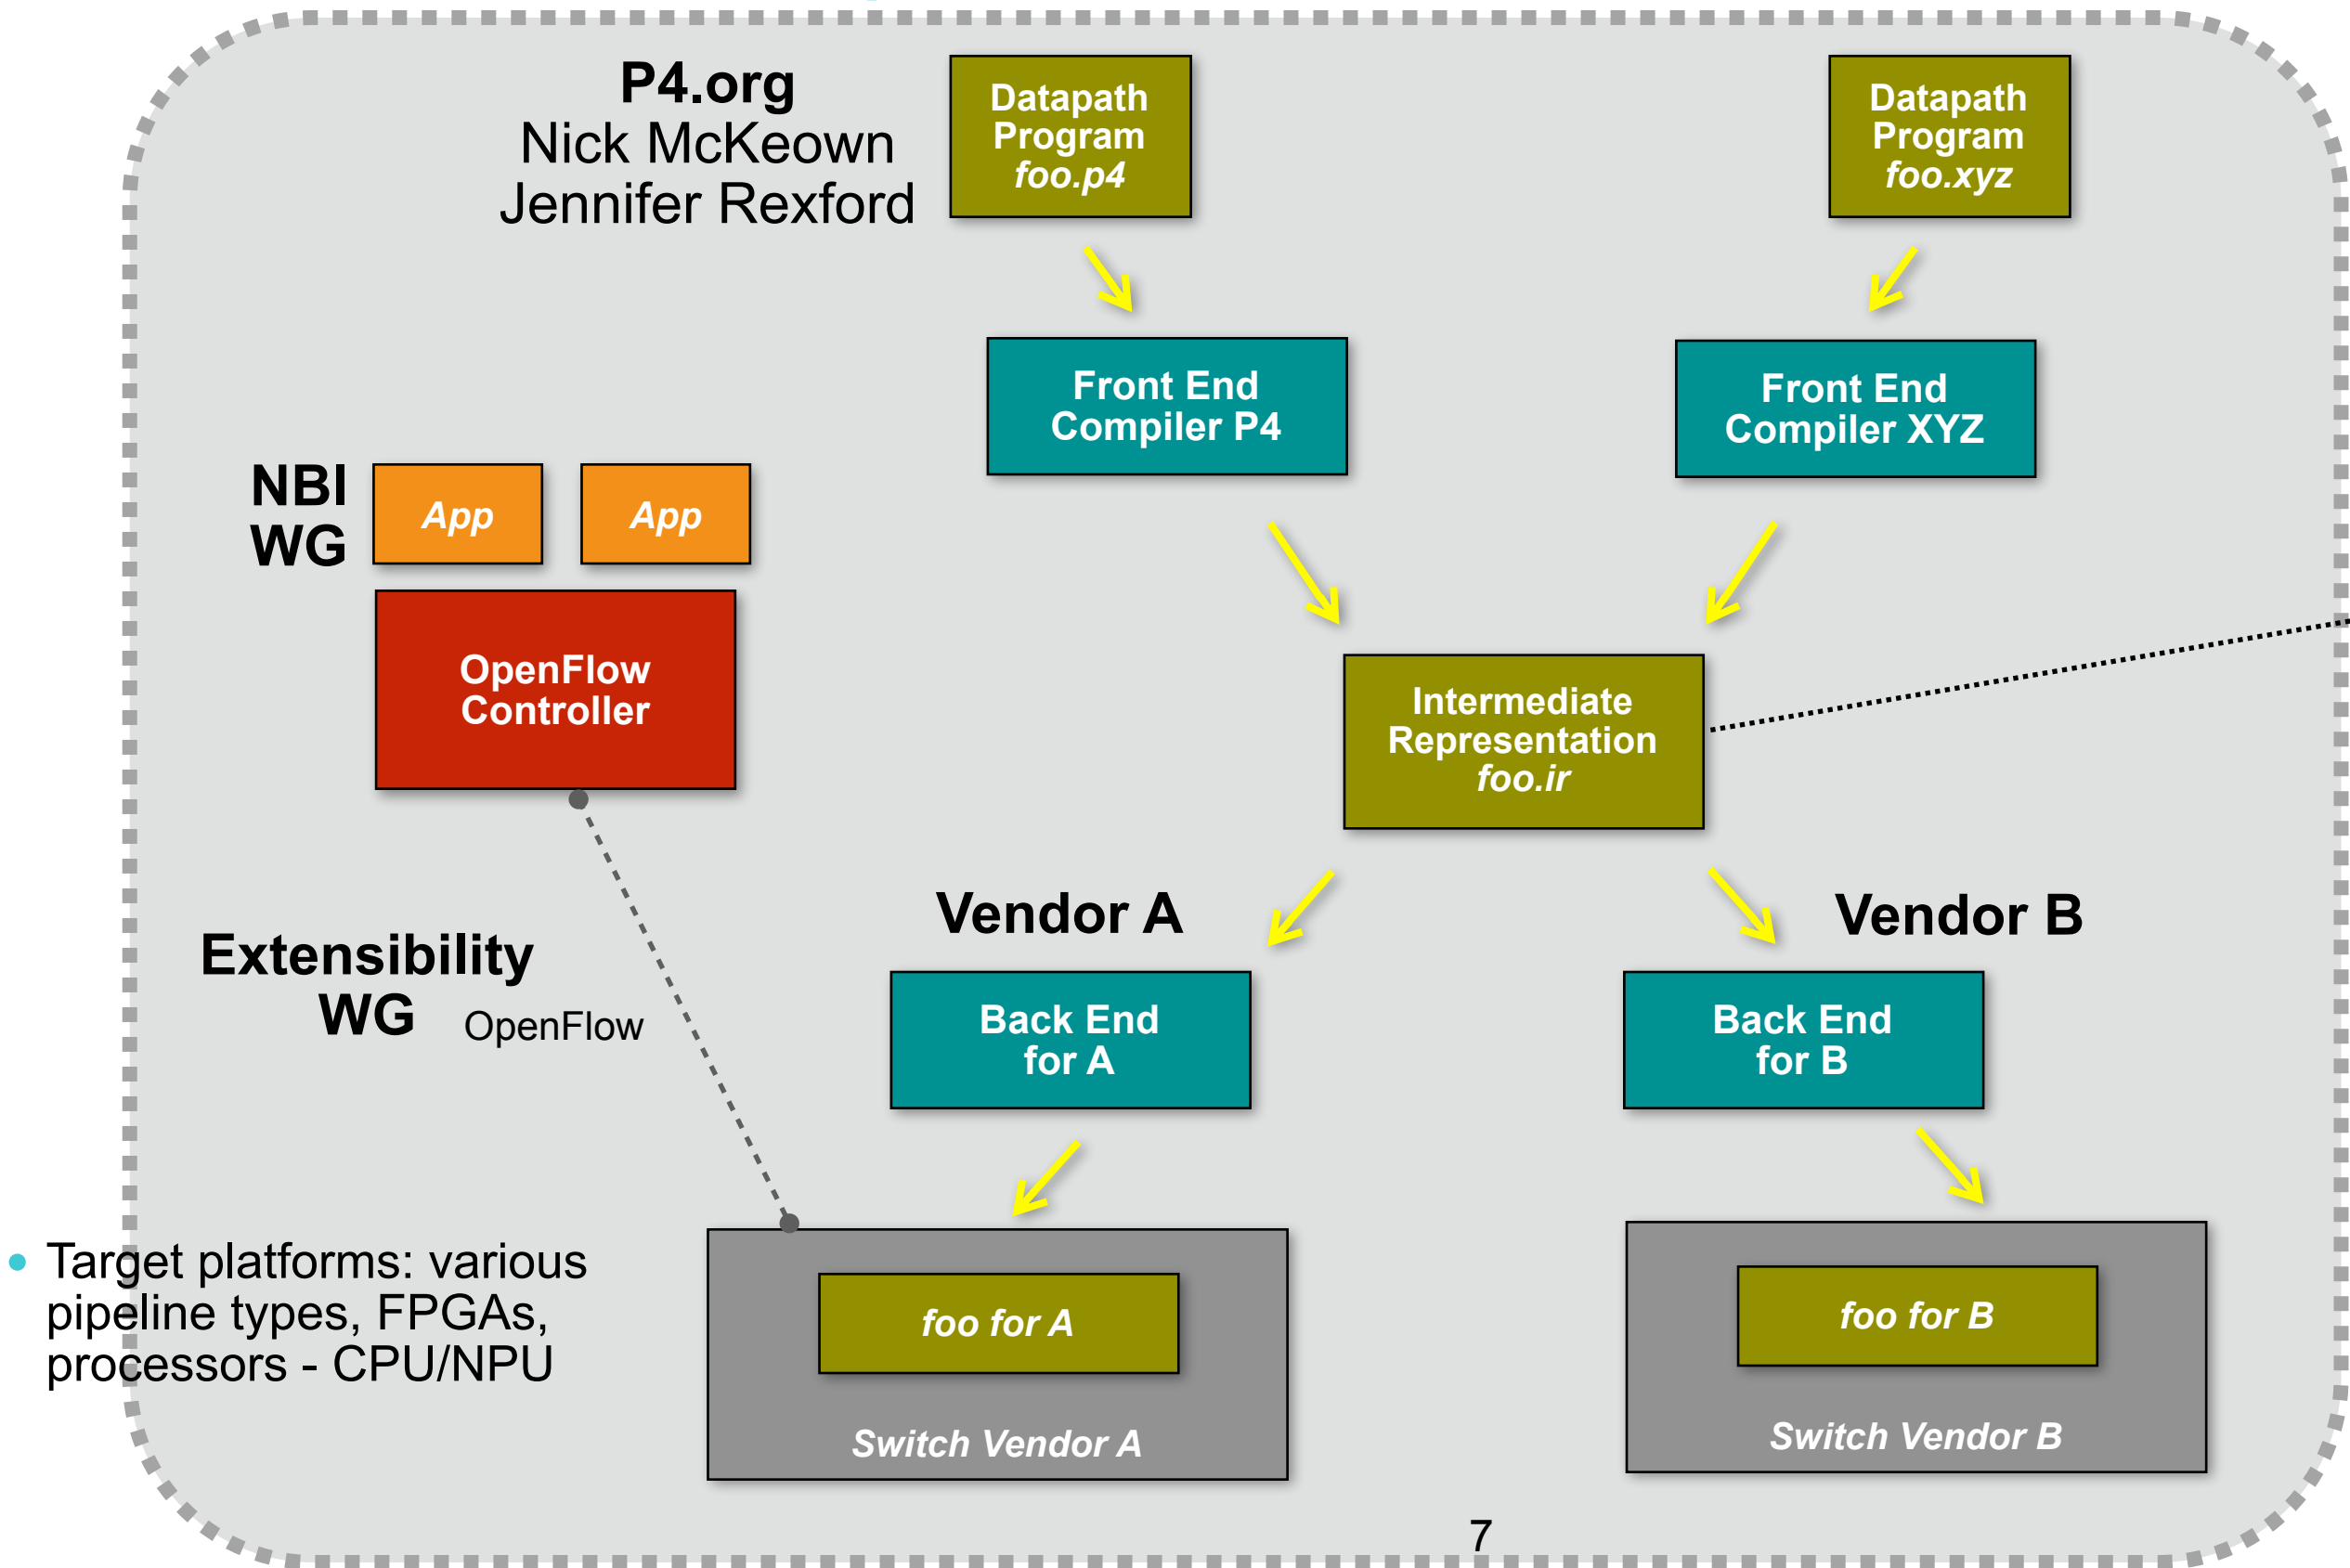
Wider OpenFlow-NG concerns e.g. lifecycles / forwarding models / capability profiles...

## ONF PIF Open Source Project

Create definition of IR by evolving implementation

## Leads to ONF Specification

FAWG? PIF Project? ArchWG? Others?



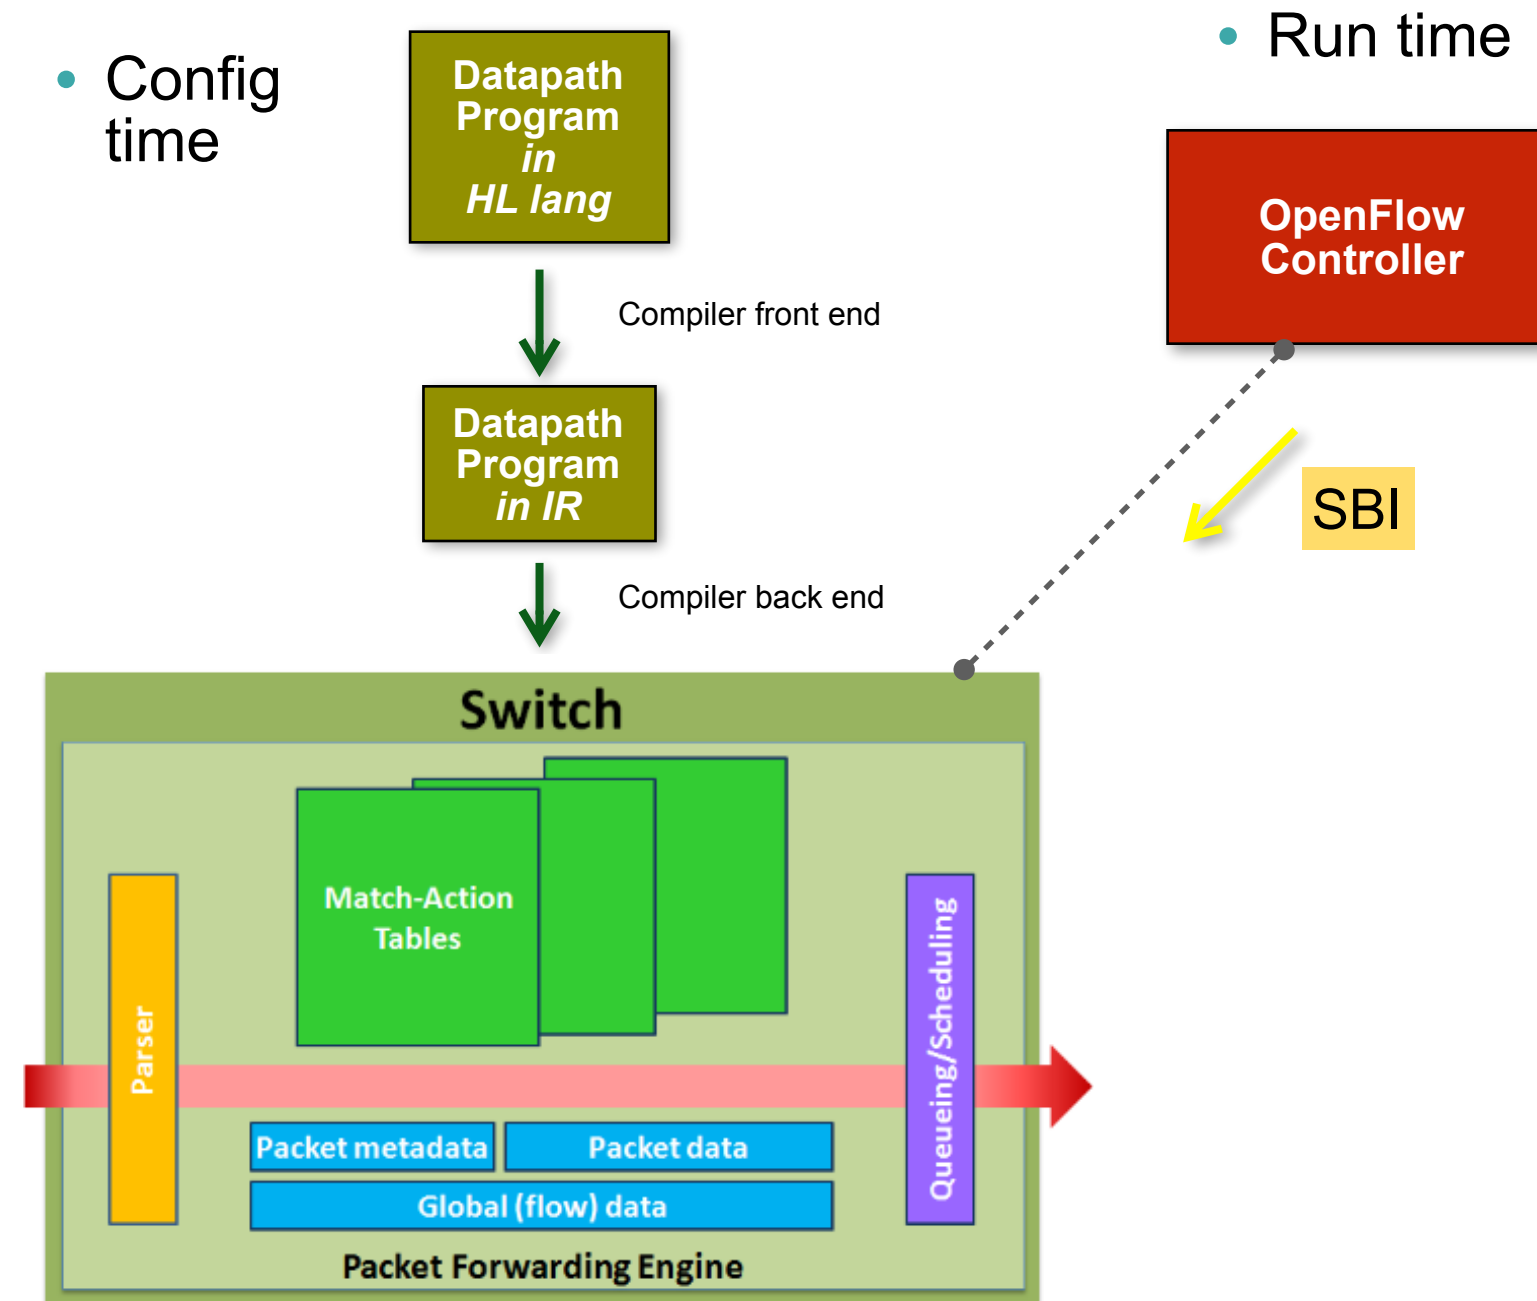

# New Elements (Draft)




- Datapath program
  - Who writes it? vendor / operator / end user / their agents
  - What is it? monolithic / modular program (libraries) => ecosystem
- Front end compiler
  - Where does this run? — developer workstation or controller?
  - Who supplies it? — multi-vendor / de-facto standard?
- Back end compiler
  - Where does this run? — switch or controller?
  - Who supplies it? — almost certainly vendor specific



# AIR-IRI: Prototype IR




# AIR-IRI: Prototype IR



- Config time

- Run time



## Now


Interpreter for candidate IR

Written in Python for Linux

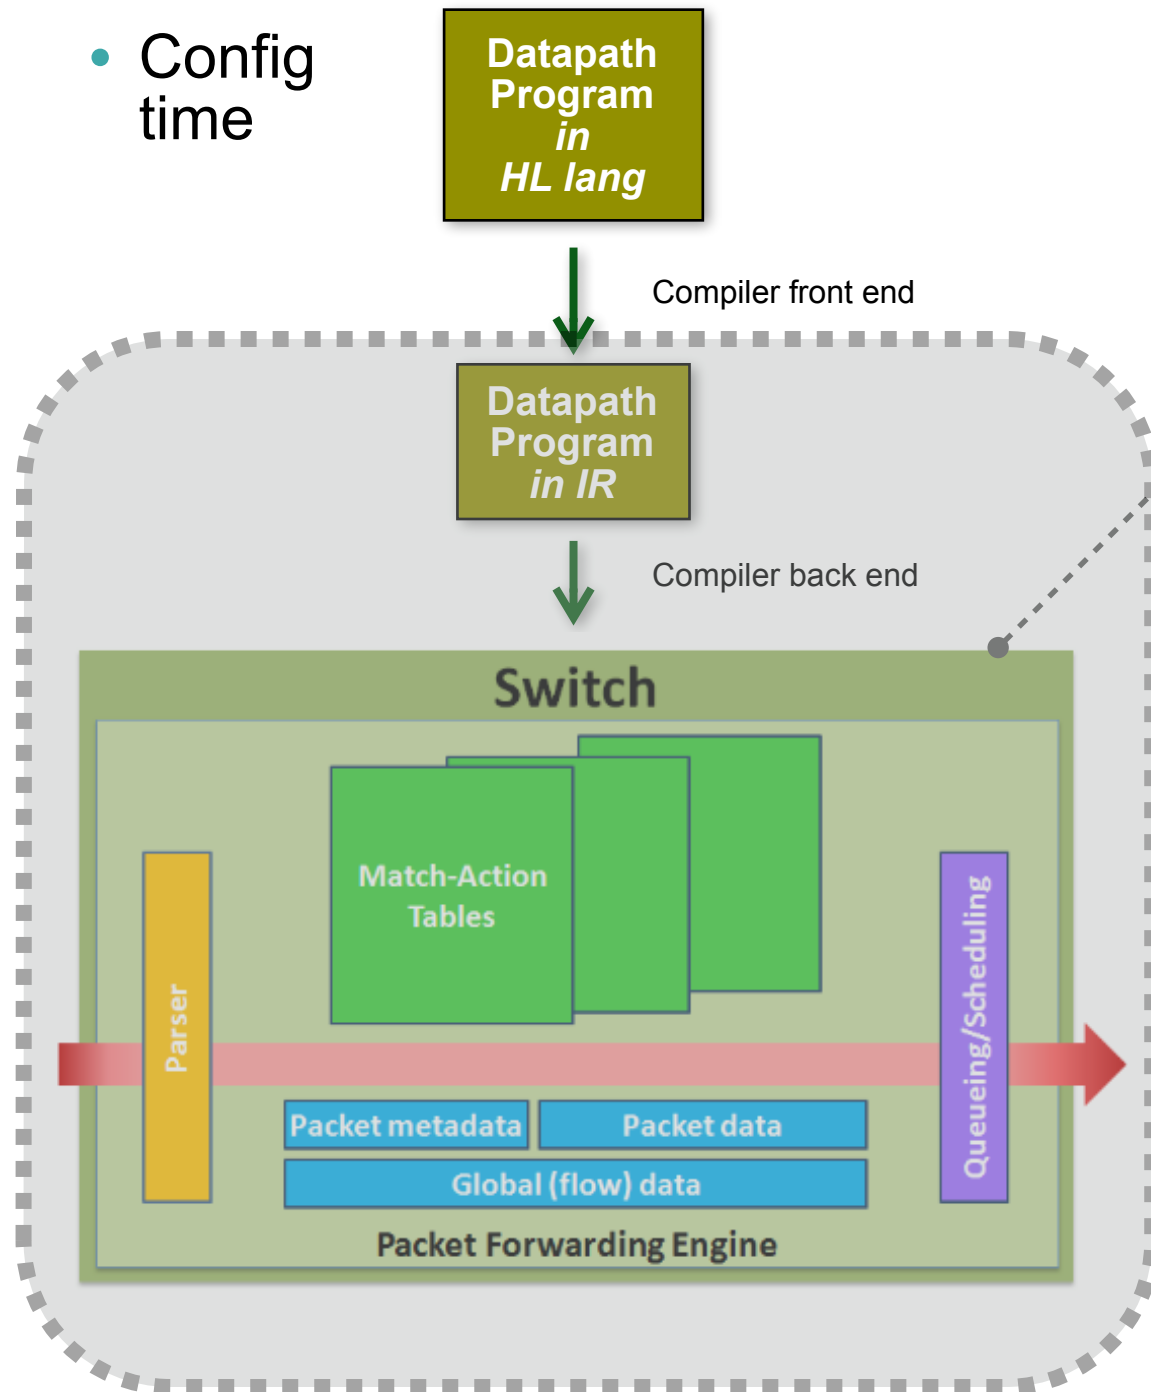
Enables evolving IR

Processes packets in PCAP files or via Linux netdevs

# AIR-IRI: Prototype IR



- Config time



**Soon**

Hook up to run-time interface  
- Apache Thrift?  
- OpenFlow 1.x?

**Now**


Interpreter for candidate IR

Written in Python for Linux


Enables evolving IR

Processes packets in PCAP files or via Linux netdevs

# AIR-IRI: Prototype IR



- Config time



## Soon

Hook up to run-time interface

- Apache Thrift?
- OpenFlow 1.x?

## Now

Interpreter for candidate IR

Written in Python for Linux

Enables evolving IR

Processes packets in PCAP files or via Linux netdevs

## ToDo for YOU!

*Define new IR variant, improve interpreter to support it*

*Implement IR back end for your platform*

# Meta IR Facilitates Evolution / Experimentation



- AIR = framework for creating IR interpreters
  - Facilitates introducing classes of “objects”, e.g. match tables, action blocks, TM blocks
  - Each “object” has named attributes (like C structure)
  - air\_meta.yml defines the objects and their attributes (declares classes/types)
- IRI = instance of an IR
  - Specific set of “objects” with implementation of each
  - Currently defined objects: value\_set, value\_map, table, header, metadata, action, parse\_state, parser, control\_flow, traffic\_manager, processor\_layout
  - Community invited to extend / modify this

# IRI Elements: Types



## # Complete list of types:

- table
- header
- metadata
- action
- parse\_state
- parser
- control\_flow
- traffic\_manager
- processor\_layout

## # Types with a process method:

- control\_flow
- parser
- traffic\_manager

# IRI Elements: Attributes per Type



```
# All support type and doc
air_attributes :
  table :
    - match_on
  header :
    - fields
    - max_depth # hdr stack if > 1
  metadata :
    - fields
    - initial_values
  action :
    - format
    - parameter_list
    - implementation
  parse_state :
    - extracts
    - select_value # Optional
```

```
# CONTINUED
control_flow : *graph_attributes
parser :
  - format
  - implementation
  - start_state
traffic_manager : # Experimental
  - queues_per_port
  - dequeue_discipline
  - egress_spec_map
processor_layout:
  - format
  - implementation
  - port_count
```

# Layout: Top Level Pipeline Structure




## **layout:**

```
type : processor_layout
doc  : "The layout specification for the switch instance"
port_count : 4
format : list
implementation :
  - parser
  - ingress_flow
  - tm_queues
  - egress_flow
```

- Changes for each category of “datapath program” (e.g. with / without QoS)
- Simple list of processors currently implemented
- Future: more complex topologies



# Protocol / Metadata Fields and Actions



```
# Header object
ethernet :
  type : header
  doc : "The L2 header"
  fields :
    - dst_mac : 48
    - src_mac : 48
    - ethertype : 16

# Metadata object
pkt_md : # General metadata
  type : metadata
  doc : "General metadata for the packet"
  fields :
    # Virtual network instance identifier
    - vni : 16

# Action object
set_vni_a :
  type : action
  doc : "Set the VNI in metadata"
  format : action_set
  parameter_list :
    - vni_id
  implementation : >-
    modify_field(pkt_md.vni, vni_id);
```

- Protocol fields vary according to targeted network protocols
- Metadata fields and actions vary according to required behavior
- =>“Datapath programming” commencing in earnest

# Protocol Details - Complete Parser



```
ethernet_p :  
  type : parse_state  
  doc : "Parse state for ethernet"  
  extracts :  
    - ethernet  
  select_value :  
    - ethernet.ethertype
```

```
vlan_p :  
  type : parse_state  
  doc : "Parse state for vlan tag"  
  extracts :  
    - vlan_tag
```

```
parser :  
  type : parser  
  doc : "Implementation of primary parser"  
  format : dot  
  start_state : ethernet_p  
  implementation : >-  
    digraph {  
      ethernet_p -> vlan_p [value="0x8100"]  
      ethernet_p -> vlan_p [value="0x9100"]  
    }
```

- Parse tree currently specified in dedicated parser object (parsing before matching)
- Being considered: match protocol ID using a table, then trigger parsing next header (parse - match - parse - match)

# Tables and Control Flow



```
vni :
  type : table
  doc : "Map VLAN to VNI"
  match_on :
    vlan_tag.vlan_id : ternary
    ethernet.src_mac : ternary


forward :
  type : table
  doc : "Forward based on L2 dest addr"
  match_on :
    pkt_md.vni : exact
    ethernet.dst_mac : exact

acl :
  type : table
  doc : "Perform ACL operations"
  match_on :
    pkt_md.vni : exact
    ethernet.dst_mac : exact
    ethernet.src_mac : exact
```

```
ingress_flow :
  type : control_flow
  doc : "The control flow for ingress"
  format : dot
  implementation : >-
    digraph {
      vni -> forward [action=set_vni_a]
      forward -> exit_control_flow [action=set_egress_a]
      forward -> exit_control_flow [action=drop_pkt_a]
    }

egress_flow :
  type : control_flow
  doc : "The control flow for egress"
  format : dot
  implementation : >-
    digraph {
      acl -> exit_control_flow [action=set_dst_mac_a]
      acl -> exit_control_flow [action=drop_pkt_a]
    }
```

# AIR / IRI Evolution



- Evolve IR language interpreter
  - Enhance existing objects e.g. matching, actions, QoS
  - New concepts e.g. statefulness
- Tools operating on IR datapath programs
  - Visualization
  - Import / Export e.g. create NDM / TTP from IR program
  - Predict performance / capacity requirements
  - Transform - e.g. parse everything then match vs. incremental distributed parsing
- Introduce run-time
  - Callable auto-marshalled, e.g. Apache Thrift
  - Traditional protocol based, e.g. OpenFlow 1.x
  - Sample controller + sample applications on controller
- Samples
  - Sample datapath applications (in IR? in high level language once front end compilers ready?)
  - Libraries for common protocols, actions (+ infrastructure for this - templating)

# Discussion + Next Steps

